

# Intro to Neural Networks II

# The Problem: Learning Slowdown I

- The *Quadratic Cost* (or Mean Squared Error) is a common starting point:

$$J = \frac{1}{2n} \sum_x \|y(x) - a^L(x)\|^2$$

- The gradient terms for weight ( $w$ ) and bias ( $b$ ) in the output layer are:

$$\frac{\partial J}{\partial w_{jk}^L} = (a_j^L - y_j) \sigma'(z_j^L) a_k^{L-1}$$

$$\frac{\partial J}{\partial b_j^L} = (a_j^L - y_j) \sigma'(z_j^L)$$

# The Problem: Learning Slowdown II

- **The Problem:** The  $\sigma'(z_j^L)$  term.
- When the neuron is "saturated" (output  $a_j^L$  is close to 0 or 1), the sigmoid function becomes very flat.
- This means  $\sigma'(z_j^L) \approx 0$ .

## Consequence

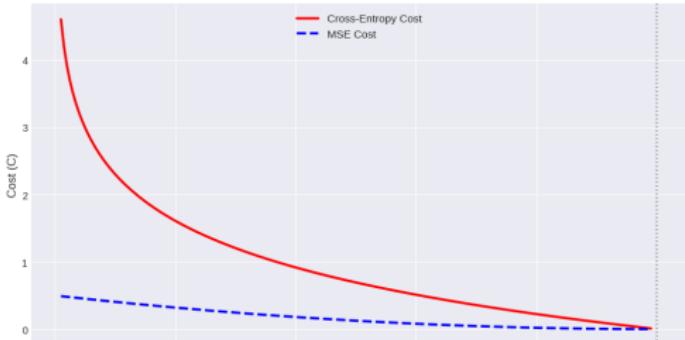
Even if the error  $(a_j^L - y_j)$  is large, the gradients become tiny, and the network learns **very slowly**.

# Visualizing Learning Slowdown

**Top:** Both are 0 when prediction is 1 (correct); both increase when prediction goes to 0 (wrong).

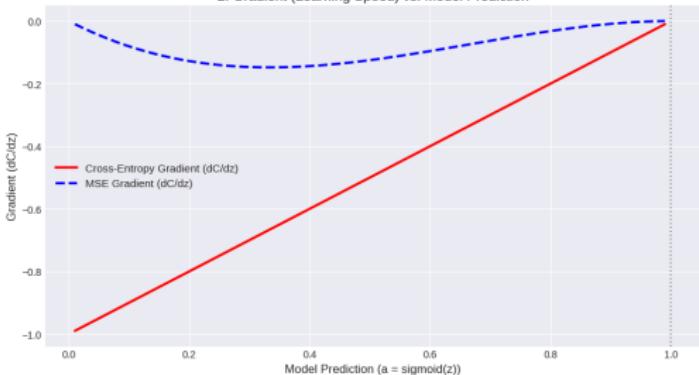
MSE vs Cross-Entropy with Sigmoid (Target y=1)

1. Cost Functions vs. Model Prediction



**Bottom:** MSE gradient goes to zero when the prediction is wrong (vanishing gradient).  
Cross-entropy gradient is zero only when prediction is 1!

2. Gradient (Learning Speed) vs. Model Prediction



# The Cross-Entropy Cost Function

## Definition

For a single neuron with one input  $x$ , output  $a = \sigma(z)$ , and target  $y$ :

$$J = -\frac{1}{n} \sum_x [y \ln(a) + (1 - y) \ln(1 - a)]$$

where  $a = \sigma(wx + b)$ .

- **Why is this a cost function?**
  - It's non-negative ( $J > 0$ ).
  - If the neuron's output  $a$  is close to the target  $y$ , the cost  $J \rightarrow 0$ . (e.g., if  $y = 0$  and  $a \rightarrow 0$ , then  $-\ln(1 - a) \rightarrow 0$ ).
- **The Big Idea:** It's designed to solve the learning slowdown problem.

# Cross-Entropy: Derivation & Properties I

- Let's calculate the gradient for the output layer weights. We need  $\frac{\partial J}{\partial w_{jk}^L}$ .
- By the chain rule:  $\frac{\partial J}{\partial w_{jk}^L} = \frac{\partial J}{\partial a_j^L} \frac{\partial a_j^L}{\partial z_j^L} \frac{\partial z_j^L}{\partial w_{jk}^L}$
- Let's compute the terms (for a single training example  $x$ ):

$$\frac{\partial J}{\partial a_j^L} = - \left( \frac{y_j}{a_j^L} - \frac{1 - y_j}{1 - a_j^L} \right) = \frac{a_j^L - y_j}{a_j^L(1 - a_j^L)}$$

$$\frac{\partial a_j^L}{\partial z_j^L} = \sigma'(z_j^L) = \sigma(z_j^L)(1 - \sigma(z_j^L)) = a_j^L(1 - a_j^L)$$

► proof

$$\frac{\partial z_j^L}{\partial w_{jk}^L} = x$$

# Cross-Entropy: Derivation & Properties II

- Combine them:

$$\frac{\partial J}{\partial w_{jk}^L} = \frac{a_j^L - y_j}{a_j^L(1 - a_j^L)} \cdot \left( a_j^L(1 - a_j^L) \right) \cdot x$$

## The Result

The  $a_j^L(1 - a_j^L)$  and  $\sigma'(z)$  terms **cancel out!**

$$\frac{\partial J}{\partial w_{jk}^L} = (a_j^L - y_j)x$$

$$\frac{\partial J}{\partial b_j^L} = (a_j^L - y_j)$$

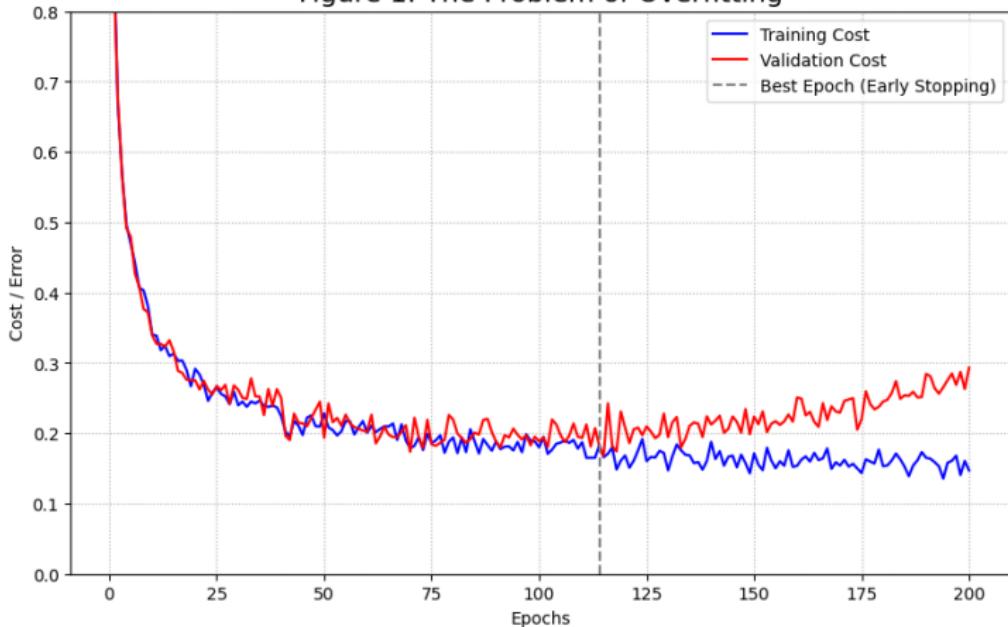
- The gradient is now proportional to the **error**  $(a_j^L - y_j)$ . A larger error means a larger gradient and faster learning.

# The Problem: Overfitting

- **Definition:** A model is **overfitting** when it learns the training data *too well*, including its noise.
- It fails to **generalize** to new, unseen data (like the validation or test sets).
- **Symptom:**
  - Training accuracy keeps improving (or training cost keeps decreasing).
  - Validation accuracy plateaus or, even worse, starts to get *worse*.
- This often happens when a model has too much capacity (too many parameters) for the amount of data available.

# Visualizing Overfitting

Figure 1: The Problem of Overfitting



# Regularization: The General Idea

- If overfitting is the problem, how do we fight it?
- **1. Get More Training Data (The Best Way)**
  - A larger, more diverse dataset is the most effective defense.
  - If this isn't possible, we can...
- **2. Artificially Augment Training Data**
  - Create "new" training samples by transforming existing ones.
  - **For images:** Rotate, translate (shift), add random noise, or apply elastic distortions.
  - This teaches the network to be robust to these variations.
- **3. Use Regularization Techniques**
  - Add a "complexity penalty" to the cost function to discourage the network from learning complex patterns.
  - We will look at L2, L1, and Dropout.

# Regularization: The General Idea

- If overfitting is the problem, how do we fight it?
- **1. Get More Training Data (The Best Way)**
  - A larger, more diverse dataset is the most effective defense.
  - If this isn't possible, we can...
- **2. Artificially Augment Training Data**
  - Create "new" training samples by transforming existing ones.
  - **For images:** Rotate, translate (shift), add random noise, or apply elastic distortions.
  - This teaches the network to be robust to these variations.
- **3. Use Regularization Techniques**
  - Add a "complexity penalty" to the cost function to discourage the network from learning complex patterns.
  - We will look at L2, L1, and Dropout.

# Regularization: The General Idea

- If overfitting is the problem, how do we fight it?
- **1. Get More Training Data (The Best Way)**
  - A larger, more diverse dataset is the most effective defense.
  - If this isn't possible, we can...
- **2. Artificially Augment Training Data**
  - Create "new" training samples by transforming existing ones.
  - **For images:** Rotate, translate (shift), add random noise, or apply elastic distortions.
  - This teaches the network to be robust to these variations.
- **3. Use Regularization Techniques**
  - Add a "complexity penalty" to the cost function to discourage the network from learning complex patterns.
  - We will look at L2, L1, and Dropout.

# Regularization: L2 (Weight Decay) I

## Idea

Penalize large weights. A network with small weights is "simpler" and less prone to overfitting.

- We add a regularization term to the cost function  $J_0$ :

$$J = J_0 + \frac{\lambda}{2n} \sum_w w^2$$

- $J_0$  = original cost (e.g., cross-entropy)
- $\lambda > 0$  = the regularization hyperparameter.

# Regularization: L2 (Weight Decay) II

- How does this change the gradient?

$$\frac{\partial J}{\partial w} = \frac{\partial J_0}{\partial w} + \frac{\lambda}{n} w$$

$$\frac{\partial J}{\partial b} = \frac{\partial J_0}{\partial b} \quad (\text{Biases are usually not regularized})$$

- The SGD Update Rule:

$$\begin{aligned}w &\rightarrow w - \eta \left( \frac{\partial J_0}{\partial w} + \frac{\lambda}{n} w \right) \\&= \left( 1 - \frac{\eta \lambda}{n} \right) w - \eta \frac{\partial J_0}{\partial w}\end{aligned}$$

# Regularization: L2 (Weight Decay) III

## Weight Decay

The term  $\left(1 - \frac{\eta\lambda}{n}\right)$  scales the weight down on every step. This is why L2 is also called **weight decay**.

# Regularization: L1 |

## Idea

Penalize the *absolute value* of the weights.

- The L1-regularized cost function:

$$J = J_0 + \frac{\lambda}{n} \sum_w |w|$$

- **The SGD Update Rule:** (where  $\text{sgn}(w)$  is -1 if  $w < 0$ , +1 if  $w > 0$ )

$$\begin{aligned} w &\rightarrow w - \eta \left( \frac{\partial J_0}{\partial w} + \frac{\lambda}{n} \text{sgn}(w) \right) \\ &= w - \frac{\eta \lambda}{n} \text{sgn}(w) - \eta \frac{\partial J_0}{\partial w} \end{aligned}$$

# Regularization: L1 ||

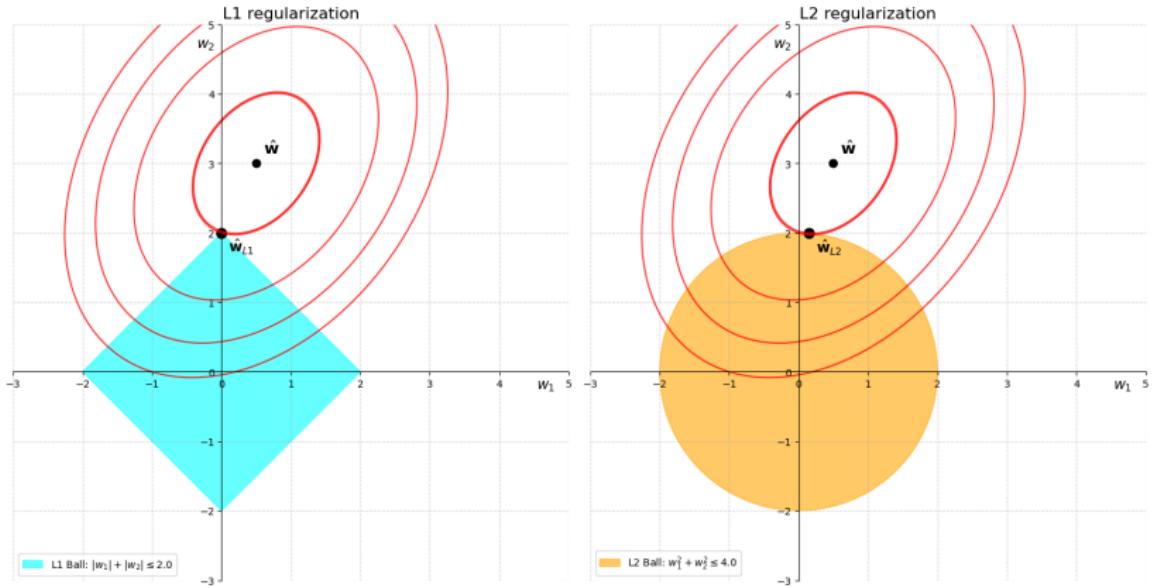
- **Difference from L2:**

- **L2 (Weight Decay):** Shrinks weights by a *proportional* amount.  
Large weights are shrunk more than small weights.
- **L1:** Shrinks weights by a *constant* amount (proportional to  $\eta\lambda/n$ ),  
pushing them towards 0.

## Sparsity

L1 regularization tends to set many weights to be **exactly zero**. This creates a **sparse model**, which can be useful for feature selection.

# Visualizing L1 vs. L2



# Regularization: Dropout I

## Idea

A radically different approach. Instead of modifying the cost, it modifies the network itself.

- **During Training (for each mini-batch):**

- ① Go through each hidden layer.
- ② For each neuron, **randomly** "drop" it (set its output to 0) with some probability  $p$  (e.g.,  $p = 0.5$ ).
- ③ Perform forward- and back-propagation on this "thinned" network.
- ④ Update weights and biases.
- ⑤ Repeat for the next mini-batch, with a new random set of dropped-out neurons.

# Regularization: Dropout II

- **During Testing / Evaluation:**
  - Use the full, complete network (no neurons are dropped).
  - **BUT:** Multiply the weights outgoing from the hidden layers by the probability  $p$  (e.g., 0.5) to compensate for the fact that more neurons are active now.
- **Intuition:** It's like training a massive **ensemble** of different, smaller networks. It forces the network to learn redundant representations, as it cannot rely on any single neuron (it might be dropped!).

# The Problem with Initialization

- **Classical Method:** Initialize weights  $w$  and biases  $b$  from a standard Gaussian (Normal) distribution with mean 0 and standard deviation 1.
- **The Problem:** Consider a neuron with  $n_{in}$  inputs.

$$z = \sum_{j=1}^{n_{in}} w_j x_j + b$$

- Assume inputs  $x_j$  are normalized (e.g.,  $\sim 50\%$  are 0,  $\sim 50\%$  are 1).
- The weighted sum  $z$  will also be a Gaussian distribution.
- **Its standard deviation will be large:**  $\text{std}(z) \approx \sqrt{n_{in} \cdot \text{std}(w)^2} = \sqrt{n_{in}}$ .

## Consequence

If  $n_{in}$  is large (e.g., 1000 inputs),  $\text{std}(z) \approx 31.6$ . This means  $z$  is very likely to be a large positive or negative number.

- This pushes the neuron's output  $\sigma(z)$  into the **saturated** regions (near 0 or 1), causing  $\sigma'(z) \approx 0$  and **slow learning** right from the start.

# Visualizing Saturation

Figure 3a: Sigmoid Activation Function

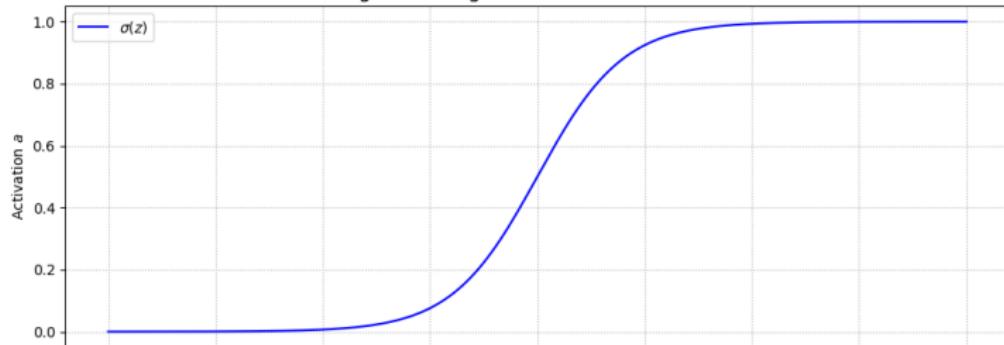
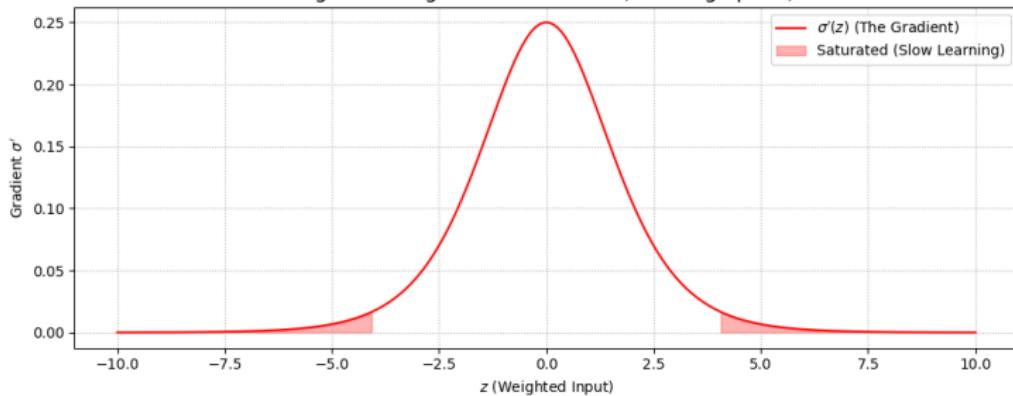


Figure 3b: Sigmoid's Derivative (Learning Speed)



# A Better Weight Initialization I

## The Solution

Initialize weights from a Gaussian distribution with mean 0, but change the standard deviation.

$$w \sim \mathcal{N}(0, 1/\sqrt{n_{in}})$$

where  $n_{in}$  is the number of input connections to the neuron.

## Why does this work?

- Let's re-calculate the standard deviation of  $z = \sum w_j x_j + b$ :

$$\text{std}(z) = \sqrt{n_{in} \cdot \text{std}(w)^2} = \sqrt{n_{in} \cdot (1/\sqrt{n_{in}})^2}$$

$$\text{std}(z) = \sqrt{n_{in} \cdot (1/n_{in})} = \sqrt{1} = 1$$

# A Better Weight Initialization II

## Benefit

The weighted sum  $z$  is now a well-behaved Gaussian  $\mathcal{N}(0, 1)$ .

- This keeps the neuron in the "active" region of the sigmoid function, where  $\sigma'(z)$  is not close to zero.
- This (along with cross-entropy) allows learning to start at a much healthier pace.
- Biases can be initialized as  $\mathcal{N}(0, 1)$  or just as 0.

# Hyperparameter Tuning as a "Black Art" I

- Hyperparameters (HPs) are the settings we don't learn, but *set* before training:
  - Learning rate  $\eta$
  - Regularization parameter  $\lambda$
  - Mini-batch size (see later)
  - Number of epochs
  - Network architecture (layers, neurons)
- There is no simple rule; it's an iterative process of experimentation.

# Hyperparameter Tuning as a "Black Art" II

## Broad Strategy

- ① **Simplify first.** Turn off all regularization ( $\lambda = 0$ , no dropout).
- ② **Find  $\eta$ .** Find the "order of magnitude" for the learning rate  $\eta$  that causes the *training cost* to start decreasing.
- ③ **How to find  $\eta$ :** Try  $\eta = 0.01$ .
  - If cost explodes/oscillates, decrease (e.g.,  $\eta = 0.001$ ).
  - If cost decreases too slowly, increase (e.g.,  $\eta = 0.1, 1.0$ ).
- ④ **Tune  $\lambda$ :** Once you have a reasonable  $\eta$ , turn on regularization. Start with  $\lambda = 1.0$  and try tuning it on a log-scale (e.g., 0.1, 10.0). Use the **validation accuracy** to pick the best  $\lambda$ .
- ⑤ **Re-tune  $\eta$ :** Your best  $\lambda$  might require a new, fine-tuned  $\eta$ .
- ⑥ **Tune Mini-batch size:** This often interacts with  $\eta$ . A smaller batch size might need a smaller  $\eta$ . Often, you pick a size that fits your GPU memory and tune  $\eta$  accordingly.

# Tuning: The Learning Rate ( $\eta$ )

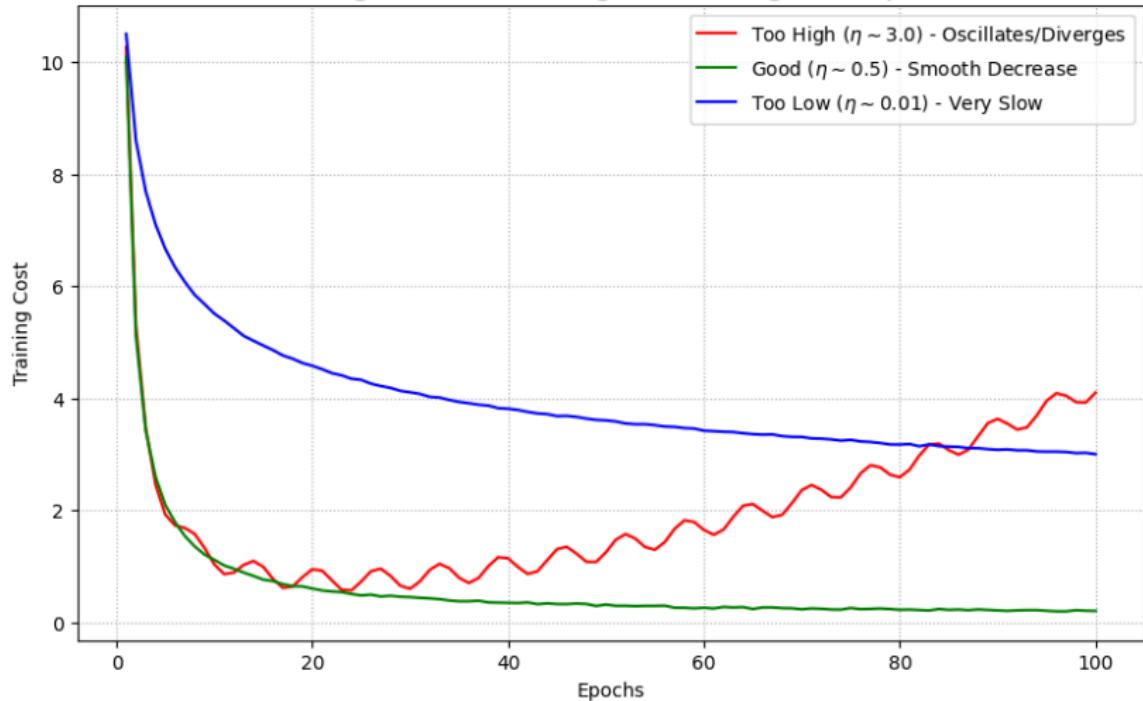
- $\eta$  is arguably the most important hyperparameter.
- **Too High:** The cost will oscillate wildly or "explode" (NaNs).
- **Too Low:** The cost will decrease, but *extremely* slowly. You'll waste time.
- **Just Right:** The cost decreases steadily.

## The "Golden Rule"

Monitor **validation accuracy**. Train with your starting  $\eta$  (e.g.,  $\eta = 0.1$ ). When the validation accuracy stops improving ("plateaus"), divide  $\eta$  by a factor (e.g., 2 or 10) and keep training. Repeat this "learning rate schedule" several times.

# Visualizing Learning Rates

Figure 4: Choosing a Learning Rate  $\eta$



# Final Summary

- **Cross-Entropy:** Solves the learning slowdown from saturated neurons by ensuring the gradient is proportional to the error.
- **Regularization:** A set of techniques (L2, L1, Dropout, Augmenting Data) to combat overfitting and help the model generalize.
- **Weight Initialization:** Using  $\mathcal{N}(0, 1/\sqrt{n_{in}})$  prevents neurons from saturating at the start of training.
- **Hyperparameters:** Tuning is an iterative process. Start simple, find a working  $\eta$ , then tune  $\lambda$  on the validation set.

## Key Takeaway

These techniques work together to make training deep networks **faster**, **more stable**, and **more effective**.

# The Sigmoid Function

## Definition

The sigmoid function,  $\sigma(x)$ , is defined as:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

## Alternative Form

For differentiation, it is helpful to write it as:

$$\sigma(x) = (1 + e^{-x})^{-1}$$

We will prove that the derivative of the sigmoid function is:

$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$

## Step 1: Apply the Chain Rule

Let  $u = 1 + e^{-x}$  and  $\sigma(u) = u^{-1}$ . The chain rule is:  $\frac{d\sigma}{dx} = \frac{d\sigma}{du} \cdot \frac{du}{dx}$

Part 1: Find  $\frac{d\sigma}{du}$

$$\frac{d\sigma}{du} = \frac{d}{du}(u^{-1})$$

$$\frac{d\sigma}{du} = -1 \cdot u^{-2} = -\frac{1}{u^2}$$

Part 2: Find  $\frac{du}{dx}$

$$\frac{d}{dx}(1 + e^{-x})$$

$$\frac{du}{dx} = -e^{-x}$$

## Step 1: Combine the Parts

- Substitute the parts:

$$\sigma'(x) = \left( -\frac{1}{u^2} \right) \cdot (-e^{-x})$$

- Substitute  $u$  back in:

$$\sigma'(x) = \left( -\frac{1}{(1 + e^{-x})^2} \right) \cdot (-e^{-x})$$

- Simplify to get the derivative:

Derivative

$$\sigma'(x) = \frac{e^{-x}}{(1 + e^{-x})^2}$$

## Step 2: Algebraic Manipulation

- Start with our derivative:

$$\sigma'(x) = \frac{e^{-x}}{(1 + e^{-x})^2}$$

- Add and subtract 1 in the numerator:

$$\sigma'(x) = \frac{(1 + e^{-x}) - 1}{(1 + e^{-x})^2}$$

- Split into two fractions:

$$\sigma'(x) = \frac{1 + e^{-x}}{(1 + e^{-x})^2} - \frac{1}{(1 + e^{-x})^2}$$

# Conclusion: The Final Form

- Simplify the expression:

$$\sigma'(x) = \frac{1}{1 + e^{-x}} - \left( \frac{1}{1 + e^{-x}} \right)^2$$

- Recall the definition:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- Substitute  $\sigma(x)$  back in:

$$\sigma'(x) = \sigma(x) - (\sigma(x))^2$$

- Factor out  $\sigma(x)$ :

Final Result

$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$

» back