

# Quasi-Newton and Levenberg-Marquardt methods

# Motivation

## Why Quasi-Newton Methods?

### The Goal

Find a tradeoff between the advantages and disadvantages of:

- **Steepest Descent:** Low cost per iteration ( $O(n)$ ), but slow (linear) convergence.
- **Newton-Raphson:** Very fast (quadratic) convergence, but high cost per iteration ( $O(n^3)$ ) and requires computing the full Hessian.

### The Quasi-Newton Aim

Develop methods that converge faster than steepest descent but have a lower operation count per iteration than Newton-Raphson.

# The Newton-Raphson Method

The Newton-Raphson method can be seen as replacing the original problem  $\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})$  with iteratively minimising a quadratic model.

## Second-Order Taylor Model

At  $\mathbf{x}_k$ , approximate  $f(\mathbf{x})$  with  $q(\mathbf{x})$ :

$$q(\mathbf{x}) = f(\mathbf{x}_k) + \langle \nabla f(\mathbf{x}_k), \mathbf{x} - \mathbf{x}_k \rangle + \frac{1}{2}(\mathbf{x} - \mathbf{x}_k)^T D^2 f(\mathbf{x}_k)(\mathbf{x} - \mathbf{x}_k)$$

If  $D^2 f(\mathbf{x}_k)$  is positive definite, the unique minimiser of  $q(\mathbf{x})$  is:

$$\mathbf{x}_{k+1} = \mathbf{x}_k - (D^2 f(\mathbf{x}_k))^{-1} \nabla f(\mathbf{x}_k)$$

This is the Newton-Raphson iterate.

# The Quasi-Newton Approach

## Approximating the Hessian

The core idea is to replace the true (and expensive) Hessian  $D^2f(\mathbf{x}_k)$  with a simpler approximation  $B_k$ .

# The Quasi-Newton Approach

## Approximating the Hessian

The core idea is to replace the true (and expensive) Hessian  $D^2f(\mathbf{x}_k)$  with a simpler approximation  $B_k$ .

## Quasi-Newton Quadratic Model

We build a new model  $p(\mathbf{x})$  using  $B_k \approx D^2f(\mathbf{x}_k)$ :

$$p(\mathbf{x}) = f(\mathbf{x}_k) + \langle \nabla f(\mathbf{x}_k), \mathbf{x} - \mathbf{x}_k \rangle + \frac{1}{2}(\mathbf{x} - \mathbf{x}_k)^T B_k (\mathbf{x} - \mathbf{x}_k)$$

If  $B_k$  is positive definite, minimising  $p(\mathbf{x})$  yields:

$$\mathbf{x}^* = \mathbf{x}_k - B_k^{-1} \nabla f(\mathbf{x}_k)$$

## From Iterate to Search Direction

Since  $B_k$  is only an approximation, the resulting step  $\mathbf{x}^* - \mathbf{x}_k$  is not a full update, but rather a search direction.

# From Iterate to Search Direction

Since  $B_k$  is only an approximation, the resulting step  $\mathbf{x}^* - \mathbf{x}_k$  is not a full update, but rather a search direction.

## Quasi-Newton Search Direction ( $\mathbf{d}_k$ )

$$\mathbf{d}_k = -B_k^{-1} \nabla f(\mathbf{x}_k)$$

This direction is then used in a line search to find the next iterate.

# From Iterate to Search Direction

Since  $B_k$  is only an approximation, the resulting step  $\mathbf{x}^* - \mathbf{x}_k$  is not a full update, but rather a search direction.

## Quasi-Newton Search Direction ( $\mathbf{d}_k$ )

$$\mathbf{d}_k = -B_k^{-1} \nabla f(\mathbf{x}_k)$$

This direction is then used in a line search to find the next iterate.

## Line Search Update

Find a suitable step length  $\alpha_k > 0$  (e.g., via Wolfe conditions that we will see later in **S3**):

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{d}_k$$

# From Iterate to Search Direction

Since  $B_k$  is only an approximation, the resulting step  $\mathbf{x}^* - \mathbf{x}_k$  is not a full update, but rather a search direction.

## Quasi-Newton Search Direction ( $\mathbf{d}_k$ )

$$\mathbf{d}_k = -B_k^{-1} \nabla f(\mathbf{x}_k)$$

This direction is then used in a line search to find the next iterate.

## Line Search Update

Find a suitable step length  $\alpha_k > 0$  (e.g., via Wolfe conditions that we will see later in **S3**):

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{d}_k$$

## Open problem

The key remaining part is: **How do we compute  $B_{k+1}$  cheaply?**

## Algorithm 1.1: Generic Quasi-Newton

### Generic Algorithm

**S0** Choose  $\mathbf{x}_0 \in \mathbb{R}^n$ , a nonsingular  $B_0 \in S^n$  (often  $B_0 = I$ ), and tolerance  $\epsilon > 0$ . Set  $k = 0$ .

## Algorithm 1.1: Generic Quasi-Newton

### Generic Algorithm

- S0** Choose  $\mathbf{x}_0 \in \mathbb{R}^n$ , a nonsingular  $B_0 \in S^n$  (often  $B_0 = I$ ), and tolerance  $\epsilon > 0$ . Set  $k = 0$ .
- S1** If  $\|\nabla f(\mathbf{x}_k)\| \leq \epsilon$  stop.

## Algorithm 1.1: Generic Quasi-Newton

### Generic Algorithm

- S0** Choose  $\mathbf{x}_0 \in \mathbb{R}^n$ , a nonsingular  $B_0 \in S^n$  (often  $B_0 = I$ ), and tolerance  $\epsilon > 0$ . Set  $k = 0$ .
- S1** If  $\|\nabla f(\mathbf{x}_k)\| \leq \epsilon$  stop.
- S2** Compute the quasi-Newton search direction:

$$\mathbf{d}_k = -B_k^{-1}\nabla f(\mathbf{x}_k)$$

## Algorithm 1.1: Generic Quasi-Newton

### Generic Algorithm

- S0** Choose  $\mathbf{x}_0 \in \mathbb{R}^n$ , a nonsingular  $B_0 \in S^n$  (often  $B_0 = I$ ), and tolerance  $\epsilon > 0$ . Set  $k = 0$ .
- S1** If  $\|\nabla f(\mathbf{x}_k)\| \leq \epsilon$  stop.
- S2** Compute the quasi-Newton search direction:

$$\mathbf{d}_k = -B_k^{-1}\nabla f(\mathbf{x}_k)$$

- S3** Perform a line search to find  $\alpha_k$  (e.g., satisfying Wolfe conditions) for  $\phi(\alpha) = f(\mathbf{x}_k + \alpha\mathbf{d}_k)$ .

## Algorithm 1.1: Generic Quasi-Newton

### Generic Algorithm

- S0** Choose  $\mathbf{x}_0 \in \mathbb{R}^n$ , a nonsingular  $B_0 \in S^n$  (often  $B_0 = I$ ), and tolerance  $\epsilon > 0$ . Set  $k = 0$ .
- S1** If  $\|\nabla f(\mathbf{x}_k)\| \leq \epsilon$  stop.
- S2** Compute the quasi-Newton search direction:

$$\mathbf{d}_k = -B_k^{-1}\nabla f(\mathbf{x}_k)$$

- S3** Perform a line search to find  $\alpha_k$  (e.g., satisfying Wolfe conditions) for  $\phi(\alpha) = f(\mathbf{x}_k + \alpha\mathbf{d}_k)$ .
- S4** Compute the new iterate:  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{d}_k$ .

## Algorithm 1.1: Generic Quasi-Newton

### Generic Algorithm

- S0** Choose  $\mathbf{x}_0 \in \mathbb{R}^n$ , a nonsingular  $B_0 \in S^n$  (often  $B_0 = I$ ), and tolerance  $\epsilon > 0$ . Set  $k = 0$ .
- S1** If  $\|\nabla f(\mathbf{x}_k)\| \leq \epsilon$  stop.
- S2** Compute the quasi-Newton search direction:

$$\mathbf{d}_k = -B_k^{-1}\nabla f(\mathbf{x}_k)$$

- S3** Perform a line search to find  $\alpha_k$  (e.g., satisfying Wolfe conditions) for  $\phi(\alpha) = f(\mathbf{x}_k + \alpha\mathbf{d}_k)$ .
- S4** Compute the new iterate:  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{d}_k$ .
- S5** Compute the new approximate Hessian  $B_{k+1}$  using a specific update rule.

## Algorithm 1.1: Generic Quasi-Newton

### Generic Algorithm

- S0** Choose  $\mathbf{x}_0 \in \mathbb{R}^n$ , a nonsingular  $B_0 \in S^n$  (often  $B_0 = I$ ), and tolerance  $\epsilon > 0$ . Set  $k = 0$ .
- S1** If  $\|\nabla f(\mathbf{x}_k)\| \leq \epsilon$  stop.
- S2** Compute the quasi-Newton search direction:

$$\mathbf{d}_k = -B_k^{-1}\nabla f(\mathbf{x}_k)$$

- S3** Perform a line search to find  $\alpha_k$  (e.g., satisfying Wolfe conditions) for  $\phi(\alpha) = f(\mathbf{x}_k + \alpha\mathbf{d}_k)$ .
- S4** Compute the new iterate:  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{d}_k$ .
- S5** Compute the new approximate Hessian  $B_{k+1}$  using a specific update rule.
- S6** Set  $k \leftarrow k + 1$  and go to S1.

# A Wish List for $B_k$ (P1, P2, P3)

What properties should our approximation  $B_k$  have?

## Properties P1-P3

**P1:**  $B_k$  should be nonsingular, so that  $\mathbf{d}_k = -B_k^{-1}\nabla f(\mathbf{x}_k)$  is well-defined.

# A Wish List for $B_k$ (P1, P2, P3)

What properties should our approximation  $B_k$  have?

## Properties P1-P3

- P1:**  $B_k$  should be nonsingular, so that  $\mathbf{d}_k = -B_k^{-1}\nabla f(\mathbf{x}_k)$  is well-defined.
- P2:**  $B_k$  should produce a descent direction, so that  $\mathbf{d}_k$  is a valid search direction (i.e.,  $\langle \nabla f(\mathbf{x}_k), \mathbf{d}_k \rangle < 0$ ).

# A Wish List for $B_k$ (P1, P2, P3)

What properties should our approximation  $B_k$  have?

## Properties P1-P3

- P1:**  $B_k$  should be nonsingular, so that  $\mathbf{d}_k = -B_k^{-1}\nabla f(\mathbf{x}_k)$  is well-defined.
- P2:**  $B_k$  should produce a descent direction, so that  $\mathbf{d}_k$  is a valid search direction (i.e.,  $\langle \nabla f(\mathbf{x}_k), \mathbf{d}_k \rangle < 0$ ).
- P3:**  $B_k$  should be symmetric, because the true Hessian is symmetric.

# The Power of (Symmetric) Positive Definiteness

Properties P1, P2, and P3 can all be satisfied by one simple requirement:

## Requirement

Require that  $B_k$  be **symmetric positive definite (SPD)**.

- If  $B_k$  is SPD, it is automatically nonsingular (P1) and symmetric (P3).
- P2 is satisfied (assuming  $\nabla f(\mathbf{x}_k) \neq \mathbf{0}$ ):

$$\langle \nabla f(\mathbf{x}_k), \mathbf{d}_k \rangle = -\nabla f(\mathbf{x}_k)^T B_k^{-1} \nabla f(\mathbf{x}_k) < 0$$

# A Wish List for $B_k$ (P4)

## Property P4

**P4:**  $B_{k+1}$  should be computable cheaply, ideally by "recycling" quantities we already have.

We define two **key vectors** using this information:

$$\delta_k = \mathbf{x}_{k+1} - \mathbf{x}_k = \alpha_k \mathbf{d}_k \quad (\text{change in position})$$

$$\gamma_k = \nabla f(\mathbf{x}_{k+1}) - \nabla f(\mathbf{x}_k) \quad (\text{change in gradient})$$

# The Secant Condition

How can  $\gamma_k$  give us info about the Hessian?

# The Secant Condition

How can  $\gamma_k$  give us info about the Hessian?

- A Taylor expansion gives:

$$\gamma_k = \nabla f(\mathbf{x}_{k+1}) - \nabla f(\mathbf{x}_k) \approx D^2 f(\mathbf{x}_k) \delta_k$$

- We enforce the condition that  $B_{k+1}$  would have perfectly predicted the gradient change  $\gamma_k$  over the step  $\delta_k$ .

# The Secant Condition

How can  $\gamma_k$  give us info about the Hessian?

- A Taylor expansion gives:

$$\gamma_k = \nabla f(\mathbf{x}_{k+1}) - \nabla f(\mathbf{x}_k) \approx D^2 f(\mathbf{x}_k) \boldsymbol{\delta}_k$$

- We enforce the condition that  $B_{k+1}$  would have perfectly predicted the gradient change  $\gamma_k$  over the step  $\boldsymbol{\delta}_k$ .

## The Secant Condition

The new approximation  $B_{k+1}$  must satisfy:

$$B_{k+1} \boldsymbol{\delta}_k = \gamma_k$$

This is the most important equation in quasi-Newton methods.

# The Secant Condition in 1D

Let's translate the N-D equation to 1D:

- The vectors  $x_k, \delta_k, \gamma_k$  all become **scalars**.
- The gradient  $\nabla f(x)$  becomes the **first derivative**  $f'(x)$ .
- The Hessian matrix  $B_{k+1}$  becomes a **scalar approximation**  $b_{k+1}$ .

# The Secant Condition in 1D

Let's translate the N-D equation to 1D:

- The vectors  $\mathbf{x}_k, \delta_k, \gamma_k$  all become **scalars**.
- The gradient  $\nabla f(x)$  becomes the **first derivative**  $f'(x)$ .
- The Hessian matrix  $B_{k+1}$  becomes a **scalar approximation**  $b_{k+1}$ .

The Secant Condition  $B_{k+1}\delta_k = \gamma_k$  becomes:

$$b_{k+1} \cdot (x_{k+1} - x_k) = f'(x_{k+1}) - f'(x_k)$$

# The Secant Condition in 1D

Let's translate the N-D equation to 1D:

- The vectors  $x_k, \delta_k, \gamma_k$  all become **scalars**.
- The gradient  $\nabla f(x)$  becomes the **first derivative**  $f'(x)$ .
- The Hessian matrix  $B_{k+1}$  becomes a **scalar approximation**  $b_{k+1}$ .

The Secant Condition  $B_{k+1}\delta_k = \gamma_k$  becomes:

$$b_{k+1} \cdot (x_{k+1} - x_k) = f'(x_{k+1}) - f'(x_k)$$

## The 1D Interpretation

Solving for our new approximation  $b_{k+1}$ , we get:

$$b_{k+1} = \frac{f'(x_{k+1}) - f'(x_k)}{x_{k+1} - x_k}$$

## Geometric Meaning of the 1D Condition

The equation  $b_{k+1} = \frac{f'(x_{k+1}) - f'(x_k)}{x_{k+1} - x_k}$  has a clear geometric meaning.

- This is the definition of a slope:  $\frac{\Delta y}{\Delta x}$ .
- The "y" values are not from the function  $f(x)$ , but from the **first derivative**,  $f'(x)$ .

# Geometric Meaning of the 1D Condition

The equation  $b_{k+1} = \frac{f'(x_{k+1}) - f'(x_k)}{x_{k+1} - x_k}$  has a clear geometric meaning.

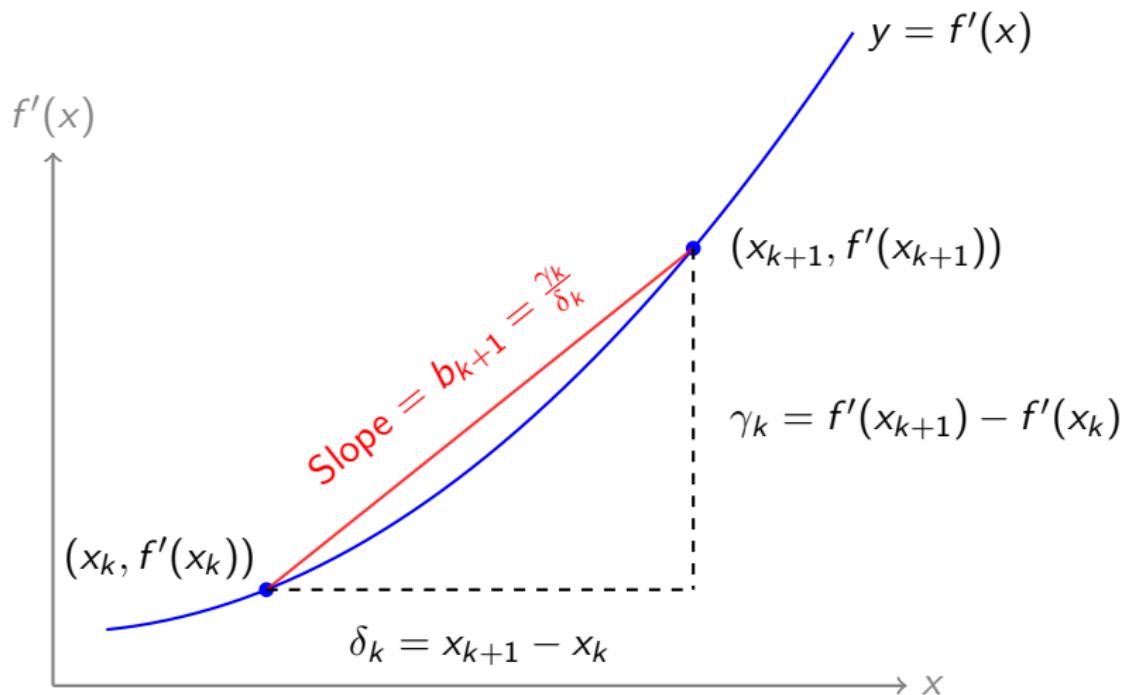
- This is the definition of a slope:  $\frac{\Delta y}{\Delta x}$ .
- The "y" values are not from the function  $f(x)$ , but from the **first derivative**,  $f'(x)$ .

## The Core Idea

The secant condition forces  $b_{k+1}$  (our approximation of the second derivative) to be the **slope of the secant line** connecting the points  $(x_k, f'(x_k))$  and  $(x_{k+1}, f'(x_{k+1}))$  on the graph of the **first derivative**.

This is a finite difference approximation of the second derivative,  $f''(x)$ .

## Geometric Meaning (Visualized)



**Figure:** The scalar  $b_{k+1}$  is the slope of the secant line on the graph of the first derivative,  $f'(x)$ .

# A Wish List for $B_k$ (P5, P6)

## Properties P5-P6

**P5:**  $B_{k+1}$  should be "close" to  $B_k$ .

- We are updating our information, not throwing it away.
- This allows  $B_k \rightarrow D^2f(\mathbf{x}^*)$  as  $\mathbf{x}_k \rightarrow \mathbf{x}^*$ .
- "Closeness" is often defined by a low-rank update.

# A Wish List for $B_k$ (P5, P6)

## Properties P5-P6

**P5:**  $B_{k+1}$  should be "close" to  $B_k$ .

- We are updating our information, not throwing it away.
- This allows  $B_k \rightarrow D^2 f(\mathbf{x}^*)$  as  $\mathbf{x}_k \rightarrow \mathbf{x}^*$ .
- "Closeness" is often defined by a low-rank update.

**P6:** The total work per iteration should be at most  $O(n^2)$ .

- This is the only way to gain a substantial speed-up over Newton-Raphson's  $O(n^3)$  cost.

# Symmetric Rank-1 (SR1) Updates I

Let's try to build an update that satisfies P3 (symmetric) and P5 (low-rank) in the simplest way possible: a symmetric rank-1 update.

## SR1 Update Form

$$B_{k+1} = B_k + \mathbf{u}\mathbf{u}^T$$

where  $\mathbf{u}$  is some vector.

- The matrix  $\mathbf{u}\mathbf{u}^T$  (outer product) is symmetric and has rank 1.
- If  $B_0$  is symmetric, all  $B_k$  will be symmetric.

# Symmetric Rank-1 (SR1) Updates II

How to find  $\mathbf{u}$ ? We enforce the Secant Condition (P4).

$$B_{k+1}\delta_k = \gamma_k \implies (B_k + \mathbf{u}\mathbf{u}^T)\delta_k = \gamma_k$$

Solving for  $\mathbf{u}$

$$(\mathbf{u}^T\delta_k)\mathbf{u} = \gamma_k - B_k\delta_k$$

i.e

$$\mathbf{u} = \frac{(\gamma_k - B_k\delta_k)}{\mathbf{u}^T\delta_k}, \quad \mathbf{u}^T = \frac{(\gamma_k - B_k\delta_k)^T}{\mathbf{u}^T\delta_k}$$

# Deriving the SR1 Formula

Substituting back into the update equation yields:

$$B_{k+1} = B_k + \frac{(\gamma_k - B_k \delta_k)(\gamma_k - B_k \delta_k)^T}{((\mathbf{u}^T \delta_k) \mathbf{u}^T) \delta_k}$$

# Deriving the SR1 Formula

Substituting back into the update equation yields:

$$B_{k+1} = B_k + \frac{(\gamma_k - B_k \delta_k)(\gamma_k - B_k \delta_k)^T}{((\mathbf{u}^T \delta_k) \mathbf{u}^T) \delta_k}$$

## SR1 Update Formula

$$B_{k+1} = B_k + \frac{(\gamma_k - B_k \delta_k)(\gamma_k - B_k \delta_k)^T}{(\gamma_k - B_k \delta_k)^T \delta_k}$$

This is the Symmetric Rank 1 (SR1) method.

# SR1: Drawbacks

The SR1 method is simple, but has serious drawbacks:

## Drawbacks

- **Positive Definiteness is not guaranteed.**  $B_{k+1}$  may not be positive definite even if  $B_k$  is. Thus,  $\mathbf{d}_k$  might not be a descent direction (P2 fails).

# SR1: Drawbacks

The SR1 method is simple, but has serious drawbacks:

## Drawbacks

- **Positive Definiteness is not guaranteed.**  $B_{k+1}$  may not be positive definite even if  $B_k$  is. Thus,  $\mathbf{d}_k$  might not be a descent direction (P2 fails).
- **The denominator can be zero or small.**

$$(\gamma_k - B_k \delta_k)^T \delta_k \approx 0$$

If this happens, the update is undefined or numerically unstable.

## SR1: The Complexity Problem (P6)

We have a problem with our  $O(n^2)$  goal (P6).

- **Computing  $B_{k+1}$ :** This involves vector operations and an outer product. This is  $O(n^2)$ . (Good!)
- **Computing  $\mathbf{d}_k$ :** To get the next search direction, we must solve  $B_k \mathbf{d}_k = -\nabla f(\mathbf{x}_k)$ . This takes  $O(n^3)$ . (Bad!)

### Problem

The  $O(n^3)$  cost of solving the system dominates the  $O(n^2)$  cost of the update, so the total iteration cost is still  $O(n^3)$ , just like Newton-Raphson.

# Solution: Sherman-Morrison-Woodbury

**Idea:** Instead of updating  $B_k$ , let's update its inverse  $H_k = B_k^{-1}$  directly.

## Theorem (Sherman-Morrison-Woodbury)

If  $B$  is  $n \times n$  and  $U, V$  are  $n \times p$ , then

$$(B + UV^T)^{-1} = B^{-1} - B^{-1}U(I_p + V^T B^{-1}U)^{-1}V^T B^{-1}$$

# Applying SMW to SR1

Suppose we know  $H_k = B_k^{-1}$  then applying the SMW formula to  $B_{k+1}$  with  $= B = B_k$ ,  $U = \mathbf{u} = (\gamma_k - B_k \delta_k)$  and  $V = U^T = \mathbf{u}^T$  we have

$$\begin{aligned} H_{k+1} &= B_{k+1}^{-1} \\ &= B_k^{-1} - B_k^{-1} \mathbf{u} (1 + \mathbf{u}^T B_k^{-1} \mathbf{u})^{-1} \mathbf{u}^T B_k^{-1} \\ &= H_k + \frac{(\delta_k - H_k \gamma_k)(\delta_k - H_k \gamma_k)^T}{(\delta_k - H_k \gamma_k)^T \gamma_k} \end{aligned} \tag{1}$$

## SR1 Inverse Update Formula

$$H_{k+1} = H_k + \frac{(\delta_k - H_k \gamma_k)(\delta_k - H_k \gamma_k)^T}{(\delta_k - H_k \gamma_k)^T \gamma_k}$$

This is also a symmetric rank-1 update, but for the inverse Hessian.

# SR1 with Inverse Hessian: P6 Solved!

Algorithm (SR1 update  $H_k$  instead of  $B_k$ )

- **S0:** Start with  $H_0 = B_0^{-1}$  (e.g.,  $H_0 = I$ ).

- **S2 (Compute  $d_k$ ):**

$$d_k = -H_k \nabla f(x_k)$$

matrix-vector multiplication  $O(n^2)$

- **S5 (Compute  $H_{k+1}$ ):**

$$H_{k+1} = H_k + \frac{(\dots)(\dots)^T}{(\dots)^T(\dots)}$$

vector additions + outer product  
 $O(n^2)$

## Complexity

The total work per iteration is now  $O(n^2)$ , satisfying P6.

## Convergence

The method converges *superlinearly* close to the minimizer.

- SR1 is a simple, rank-1 update that satisfies the secant condition.
- By updating the inverse Hessian  $H_k$ , it achieves an  $O(n^2)$  iteration cost.

## Problem

The SR1 update for  $H_k$  still has issues:

- The denominator  $(\delta_k - H_k \gamma_k)^T \gamma_k$  can be zero.
- $H_{k+1}$  (and  $B_{k+1}$ ) are not guaranteed to be positive definite.

# Motivation for BFGS

- The SR1 method is good, but the lack of guaranteed positive definiteness is a serious practical problem.

# Motivation for BFGS

- The SR1 method is good, but the lack of guaranteed positive definiteness is a serious practical problem.
- **Question:** Can we find an update that satisfies all 6 properties, especially P1-P6?

# Motivation for BFGS

- The SR1 method is good, but the lack of guaranteed positive definiteness is a serious practical problem.
- **Question:** Can we find an update that satisfies all 6 properties, especially P1-P6?
- **Answer:** Yes, but we need a slightly more complex update.

## BFGS (Broyden-Fletcher-Goldfarb-Shanno)

- The most widely used and effective quasi-Newton algorithm.
- It is a symmetric rank-2 update ( $B_{k+1} = B_k + \mathbf{u}\mathbf{u}^T + \mathbf{v}\mathbf{v}^T$ ).
- It overcomes the weaknesses of SR1 while retaining its  $O(n^2)$  speed.

# The BFGS Formula (for $B_k$ )

The BFGS update  $B_k \rightarrow B_{k+1}$  that maintains symmetry, is "close" to  $B_k$ , and satisfies the secant condition is:

## BFGS Update Formula

$$B_{k+1} = B_k - \frac{B_k \delta_k \delta_k^T B_k}{\delta_k^T B_k \delta_k} + \frac{\gamma_k \gamma_k^T}{\gamma_k^T \delta_k}$$

- This is a rank-2 update.
- Again, computing  $B_{k+1}$  is  $O(n^2)$ , but solving  $B_k d_k = -\nabla f(x_k)$  is  $O(n^3)$ .
- We could use SMW, but there is a more stable way to maintain positive definiteness.

# BFGS Strategy: Cholesky Factorisation

**Idea:** Instead of updating  $B_k$  or  $H_k$ , we will update the **Cholesky factor**  $L_k$  of  $B_k$ .

## Definition: Cholesky Factorisation

A symmetric matrix  $B$  has a Cholesky factorisation if there exists a **lower-triangular** matrix  $L$  with **positive diagonal** entries such that:

$$B = LL^T$$

## Proposition: Existence

A matrix  $B$  has a unique Cholesky factorisation  $B = LL^T$  **if and only if**  $B$  is symmetric positive definite.

This is perfect! If we can maintain an update  $L_k \rightarrow L_{k+1}$ , we automatically guarantee  $B_{k+1} = L_{k+1}L_{k+1}^T$  is positive definite (P1, P2, P3).

## Solving with Cholesky: P6 Solved!

If we have the Cholesky factor  $L_k$ , we can solve the  $O(n^3)$  system  $B_k \mathbf{d}_k = -\nabla f(x_k)$  in  $O(n^2)$  time.

Substitute  $B_k = L_k L_k^T$ : 
$$L_k L_k^T \mathbf{d}_k = -\nabla f(x_k)$$

We solve this in two (cheap) steps:

**① Forward substitution:** Let  $\mathbf{g}_k = L_k^T \mathbf{d}_k$ . Solve for  $\mathbf{g}_k$ :

$$L_k \mathbf{g}_k = -\nabla f(x_k)$$

(This is  $O(n^2)$  because  $L_k$  is lower-triangular.)

**② Backward substitution:** Solve for  $\mathbf{d}_k$ :

$$L_k^T \mathbf{d}_k = \mathbf{g}_k$$

(This is  $O(n^2)$  because  $L_k^T$  is upper-triangular.)

### Result

The search direction  $\mathbf{d}_k$  is found in  $O(n^2)$  operations.

# BFGS Derivation: High-Level Idea

How to find  $L_{k+1}$  from  $L_k$ ? (Notation:  $B = B_k$ ,  $L = L_k$ ,  $B_+ = B_{k+1}$ , etc.)

- ① We need  $B_+ = L_+ L_+^T$  to satisfy the secant condition  $B_+ \delta = \gamma$ .
- ② We also want  $L_+$  to be "close" to  $L$ .

## The Derivation

- ① Define a general (non-triangular) factor  $J$  such that  $B = JJ^T$ .
- ② Find a new factor  $J_+$  that is closest to  $J$  in Frobenius norm:

$$\min_{J_+} \|J_+ - J\|_F \quad \text{s.t.} \quad B_+ \delta = \gamma$$

(This is tricky, so a related problem is solved).

- ③ The problem solved is:  $\min_{J_+} \|J_+ - J\|_F$  s.t.  $J_+ \mathbf{g} = \gamma$  for a specific vector  $\mathbf{g}$ .

# BFGS Derivation: Step 1 (Solving for $J_+$ ) (optional)

The problem is to find the closest point  $J_+$  to  $J$  on an affine subspace.

## Minimisation Problem

$$\min_{J_+} \text{tr}((J_+ - J)(J_+ - J)^T) \quad \text{s.t.} \quad J_+ \mathbf{g} = \gamma$$

The unique solution  $J_+^*$  is found when  $(J_+^* - J)$  is orthogonal to the constraint subspace, which gives:

## Solution for $J_+$

$$J_+^* = J + \frac{(\gamma - J\mathbf{g})\mathbf{g}^T}{\mathbf{g}^T \mathbf{g}}$$

This  $J_+$  is a rank-1 update to  $J$ . Now we must choose the vector  $\mathbf{g}$ .

## BFGS Derivation: Step 2 (Finding $g$ ) (optional)

A second condition,  $J_+^T \delta = g$ , is imposed to satisfy the secant condition.

- This implies that  $g$  must be proportional to  $J^T \delta$ .

$$g = \beta J^T \delta \quad (\text{for some scalar } \beta \neq 0)$$

- Substituting this back into the relation for  $g$  and solving for  $\beta$  yields:

$$\beta = \pm \sqrt{\frac{\gamma^T \delta}{\delta^T B \delta}}$$

- **Crucial Check:** Is the square root real?

- From the Wolfe conditions, we know  $\gamma^T \delta > 0$ .
- Since  $B_k$  is positive definite,  $\delta^T B_k \delta > 0$ .
- Yes,  $\beta$  is a real, non-zero number.

## BFGS Derivation: Step 3 (The $B_+$ Update) (optional) |

Now we form the new Hessian approximation  $B_+ = J_+^* (J_+^*)^T$ .

$$B_+ = \left( J + \frac{(\dots)\mathbf{g}^T}{\mathbf{g}^T \mathbf{g}} \right) \left( J + \frac{(\dots)\mathbf{g}^T}{\mathbf{g}^T \mathbf{g}} \right)^T$$

After substituting  $\mathbf{g} = \beta J^T \boldsymbol{\delta}$  and a lot of algebra, this simplifies to:

$$B_+ = B - \frac{B\boldsymbol{\delta}\boldsymbol{\delta}^T B}{\boldsymbol{\delta}^T B\boldsymbol{\delta}} + \frac{\boldsymbol{\gamma}\boldsymbol{\gamma}^T}{\beta^2 \boldsymbol{\delta}^T B\boldsymbol{\delta}}$$

Finally, substituting  $\beta^2 = \frac{\boldsymbol{\gamma}^T \boldsymbol{\delta}}{\boldsymbol{\delta}^T B \boldsymbol{\delta}}$ :

### BFGS Formula (re-derived)

$$B_+ = B - \frac{B\boldsymbol{\delta}\boldsymbol{\delta}^T B}{\boldsymbol{\delta}^T B\boldsymbol{\delta}} + \frac{\boldsymbol{\gamma}\boldsymbol{\gamma}^T}{\boldsymbol{\gamma}^T \boldsymbol{\delta}}$$

### Key Property

This final formula for  $B_+$  is independent of  $\beta$  and, most importantly,  
**independent of the choice of  $J$ !**

## Updating the Cholesky Factor $L_k$ (optional)

- Since the  $B_+$  formula is independent of  $J$ , we can choose  $J = L_k$ , our lower-triangular Cholesky factor.
- The update formula for  $J_+$  was:

$$J_+ = J + \frac{(\gamma - \beta B\delta)\delta^T J}{\beta\delta^T B\delta}$$

- Taking the transpose gives:

$$J_+^T = J^T + \frac{J^T \delta (\gamma - \beta B\delta)^T}{\beta\delta^T B\delta} = L_k^T + (\text{rank-1 update})$$

### Problem:

$L_k^T$  is upper-triangular, but adding a rank-1 matrix makes  $J_+^T$  not triangular.

# The QR Factorisation Trick (optional)

We have a matrix  $J_+^T$  that is "almost" upper-triangular (it's  $R + uv^T$ ). How do we get our new Cholesky factor  $L_+$  from this?

## Proposition: QR Factorisation

Any square matrix  $X$  has a unique factorisation  $X = QR$ , where  $Q$  is orthogonal ( $Q^T Q = I$ ) and  $R$  is upper-triangular with non-negative diagonal.

- ① We have our non-triangular  $J_+^T$ .
- ② Compute its QR factorisation:  $J_+^T = Q_+ R_+$ .
- ③ This can be done efficiently in  $O(n^2)$  time because  $J_+^T$  is almost triangular (using Givens rotations).
- ④ Now look at  $B_+$ :

$$B_+ = J_+ J_+^T = (J_+^T)^T J_+^T = (Q_+ R_+)^T (Q_+ R_+)$$

$$B_+ = R_+^T Q_+^T Q_+ R_+ = R_+^T I R_+ = R_+^T R_+$$

- ⑤ **Result:**  $B_+ = (R_+^T)(R_+^T)^T$ .

## Algorithm: The BFGS Method

S0 Choose  $x_0$ ,  $L_0$  (e.g.,  $L_0 = I$ ), and  $\epsilon > 0$ . Set  $k = 0$ .

S1 If  $\|\nabla f(x_k)\| \leq \epsilon$ , **stop**.

S2 (**Solve for  $d_k$  in  $O(n^2)$** )

- Solve  $L_k g_k = -\nabla f(x_k)$  for  $g_k$  (forward-sub).
- Solve  $L_k^T d_k = g_k$  for  $d_k$  (back-sub).

S3 Perform line search for  $\alpha_k > 0$  satisfying Wolfe conditions.

S4 Set  $\delta_k = \alpha_k d_k$ ,  $x_{k+1} = x_k + \delta_k$ . Compute  $\gamma_k = \nabla f(x_{k+1}) - \nabla f(x_k)$ .

S5 (**Update  $L_k$  in  $O(n^2)$** )

- Compute  $J_{k+1}^T = L_k^T + (\text{rank-1 update vector})$ .
- Compute QR factorisation:  $J_{k+1}^T = Q_{k+1} R_{k+1}$ .
- Set  $L_{k+1} = R_{k+1}^T$ .

S6 Set  $k \leftarrow k + 1$  and go to S1.

## BFGS Properties

- The BFGS algorithm spends only  $O(n^2)$  computation time per iteration (P6).
- By updating the Cholesky factor  $L_k$ , it guarantees that every  $B_k$  is symmetric positive definite (P1, P2, P3).
- It uses the secant condition (P4) and is a low-rank update (P5).
- It has local **superlinear** convergence (not quite quadratic, but very fast).
- On a strictly convex quadratic function, it finds the exact minimum in at most  $n$  iterations (with exact line search).

# Final Comments & Summary Table

Let  $\mathcal{C}(f)$  be the cost of one function evaluation, then we have the following table:

Cost-Convergence Summary

Method	Cost per Iteration	Convergence Rate
Steepest Descent	$O(n\mathcal{C}(f))$	Linear
Quasi-Newton	$O(n^2 + n\mathcal{C}(f))$	Superlinear
Newton-Raphson	$O(n^3 + n^2\mathcal{C}(f))$	Quadratic

# The Line Search Problem (S3)

## The 1D Problem

Given a starting point  $\mathbf{x}_k$  and a descent direction  $\mathbf{d}_k$ , solve:

$$\min_{\alpha > 0} \phi(\alpha) \quad \text{where} \quad \phi(\alpha) = f(\mathbf{x}_k + \alpha \mathbf{d}_k)$$

## Why "Practical" Line Search?

Finding the exact minimiser  $\alpha^*$  is too computationally expensive. We just need a "good enough"  $\alpha_k$ .

# What is a "Good" Step Length?

Our  $\alpha_k$  needs to balance two competing goals:

- ① **Goal 1: Sufficient Decrease.** Avoid tiny, trivial steps.
- ② **Goal 2: Sufficient Progress.** Step shouldn't be too small to avoid getting stuck.

## The Problem

Just requiring  $f(\mathbf{x}_{k+1}) < f(\mathbf{x}_k)$  is not enough to guarantee convergence.  
The Wolfe conditions provide a "reasonable spot".

# The Wolfe Conditions

We need  $\alpha_k > 0$  and  $0 < c_1 < c_2 < 1$ .

## Condition 1: Sufficient Decrease (Armijo Rule)

$$f(\mathbf{x}_k + \alpha_k \mathbf{d}_k) \leq f(\mathbf{x}_k) + c_1 \alpha_k \nabla f(\mathbf{x}_k)^T \mathbf{d}_k$$

Rules out steps that are "too long".

## Condition 2: Curvature Condition

$$\nabla f(\mathbf{x}_k + \alpha_k \mathbf{d}_k)^T \mathbf{d}_k \geq c_2 \nabla f(\mathbf{x}_k)^T \mathbf{d}_k$$

Rules out steps that are "too short".

- Since  $\phi'(0)$  is negative,  $c_2\phi'(0)$  is a *less negative* number (e.g.,  $-0.9 \times$  slope).
- We require the new slope to be greater (less steep) than this.
- This condition rules out  $\alpha$  values that are "too short".

# Practical Line Search Algorithms I

## Backtracking Line Search

A simple and popular method that often only enforces the first condition (Armijo):

- ① Start with a full step:  $\alpha = 1$  (the "ideal" QN/Newton step).
- ② Choose  $\rho \in (0, 1)$  (e.g.,  $\rho = 0.5$ ) and  $c_1$  (e.g.,  $c_1 = 10^{-4}$ ).
- ③ **while**  $f(x_k + \alpha d_k) > f(x_k) + c_1 \alpha \nabla f(x_k)^T d_k$ :
- ④      $\alpha \leftarrow \rho \alpha$  (reduce step size)
- ⑤ **end while**
- ⑥ Set  $\alpha_k = \alpha$ .

## Full Wolfe Conditions

- More complex algorithms (using bisection and interpolation) can find an  $\alpha_k$  that satisfies *both* conditions.
- Enforcing both conditions is crucial for proving the global convergence of algorithms like BFGS.

## Drawbacks of BFGS: The Scaling Issue

While BFGS has excellent convergence properties, it faces significant challenges as the dimension  $n$  grows large (e.g., in Machine Learning or deep learning).

### The Memory Bottleneck

BFGS requires storing the approximate Hessian  $B_k$  (or its inverse/Cholesky factors) as a dense  $n \times n$  matrix.

- **Memory Complexity:**  $O(n^2)$ .

### Practical Example

Consider an optimization problem with  $n = 100,000$  variables:

- The matrix contains  $10^{10}$  entries.
- Using double precision (8 bytes per entry), this requires  $\approx 80$  GB of RAM.
- For  $n = 1,000,000$ , this becomes prohibitive ( $\approx 8$  TB).

# Introducing L-BFGS (Limited-Memory BFGS)

To tackle large-scale problems, we need a method that approximates Newton steps without the massive memory footprint.

## The Core Idea

Instead of storing the full matrix  $H_k$ , we implicitly store the approximation using a history of recent updates.

## Storage Strategy

Recall that the BFGS update is constructed using the curvature pairs:

$$\delta_k = \mathbf{x}_{k+1} - \mathbf{x}_k \quad \text{and} \quad \gamma_k = \nabla f(\mathbf{x}_{k+1}) - \nabla f(\mathbf{x}_k)$$

L-BFGS stores only the **most recent  $m$  pairs**  $\{(\delta_i, \gamma_i)\}$  (where  $m$  is small, e.g.,  $m \approx 10$  to 20). Older information is discarded.

# How L-BFGS Works

## Implicit Multiplication

We never explicitly form the matrix  $H_k$ .

- To compute the search direction  $\mathbf{d}_k = -H_k \nabla f(\mathbf{x}_k)$ , we use a specialized "**Two-Loop Recursion**" algorithm.
- This algorithm performs the matrix-vector product using only the stored pairs  $\{(\delta_i, \gamma_i)\}$  and the current gradient.

# L-BFGS: The Two-Loop Recursion

**Goal:** Compute the search direction  $\mathbf{d}_k = -H_k \nabla f(\mathbf{x}_k)$  efficiently.

## The Challenge

We cannot compute  $\mathbf{d}_k$  directly because we **do not store** the  $n \times n$  matrix  $H_k$ . We only store  $m$  vector pairs  $\{(\delta_i, \gamma_i)\}$ .

## The Core Idea

The matrix  $H_k$  is *implicitly* defined by applying the BFGS inverse update formula  $m$  times to an initial "seed" matrix  $H_k^0$  (e.g.,  $H_k^0 = I$ ).

The inverse update formula is:

$$H_{i+1} = (I - \rho_i \delta_i \gamma_i^T) H_i (I - \rho_i \gamma_i \delta_i^T) + \rho_i \delta_i \delta_i^T$$

where  $\rho_i = 1/(\gamma_i^T \delta_i)$ .

The two-loop algorithm "unrolls" this recursion to compute the matrix-vector product  $H_k \nabla f(\mathbf{x}_k)$  without ever forming  $H_k$ .

# L-BFGS: The Algorithm I

This computes  $\mathbf{r} = H_k \mathbf{q}$ , where  $\mathbf{q} = \nabla f(\mathbf{x}_k)$ .

**Input:** Gradient  $\mathbf{q} = \nabla f(\mathbf{x}_k)$ , history  $\{(\delta_i, \gamma_i)\}_{i=k-m}^{k-1}$ .

## Loop 1: Backward Pass

(Iterates from newest  $k - 1$  to oldest  $k - m$ )

**for**  $i = k - 1$  **down to**  $k - m$

$$\rho_i = 1/(\gamma_i^T \delta_i)$$

$$\alpha_i = \rho_i (\delta_i^T \mathbf{q}) \quad (*\text{Store this scalar}*)$$

$$\mathbf{q} = \mathbf{q} - \alpha_i \gamma_i$$

## Middle Step: Apply Initial Hessian

$$\mathbf{r} = H_k^0 \mathbf{q} \quad (*\text{e.g., } \mathbf{r} = \mathbf{q} \text{ if } H_k^0 = I*)$$

# L-BFGS: The Algorithm II

## Loop 2: Forward Pass

(Iterates from oldest  $k - m$  to newest  $k - 1$ )

**for**  $i = k - m$  **up to**  $k - 1$

$$\beta = \rho_i(\gamma_i^T \mathbf{r})$$

$$\mathbf{r} = \mathbf{r} + (\alpha_i - \beta)\delta_i \quad (*\text{Use stored } \alpha_i*)$$

**Output:** The search direction is  $\mathbf{d}_k = -\mathbf{r}$ .

# L-BFGS: Why It Works I

## A Nested Expression

Multiplying  $H_k \mathbf{q}$  is like evaluating a nested expression from the inside out.

$$H_k \approx (\dots (V_{k-1}^T H_k^0 V_{k-1} + \dots) \dots)$$

Where  $V_i = (I - \rho_i \gamma_i \delta_i^T)$  and we also add the rank-1 terms.

## Visualizing the Loops

- **Loop 1 (Backward Pass):** This pass moves from the "outside" of the nested expression to the "inside". It recursively applies the  $V_i$  terms (right-multiplies) to the gradient  $\mathbf{q}$  until it reaches the  $H_k^0$  core.
- **Middle Step:** The "core" matrix  $H_k^0$  is applied.
- **Loop 2 (Forward Pass):** This pass moves back "outward". It applies the  $V_i^T$  terms (left-multiplies) and adds the rank-1  $\delta_i \delta_i^T$  components, using the  $\alpha_i$  values saved from the first loop.

## The Payoff

This entire operation uses only vector dot products and additions.

- No  $n \times n$  matrices are ever formed or stored.
- **Computation Cost:**  $O(mn)$
- **Storage Cost:**  $O(mn)$

# BFGS vs L-BFGS

## Comparison: BFGS vs. L-BFGS

Feature	Standard BFGS	L-BFGS
Memory Usage	$O(n^2)$	$O(mn)$ (Linear!)
CPU per Iteration	$O(n^2)$	$O(mn)$
Convergence	Superlinear	Linear (but fast)
Use Case	Small/Medium $n$	Large scale $n > 1000$

# Non-Linear LS: the Levenberg-Marquardt algorithm.

## Goal

We want to find the parameters  $\mathbf{w}$  that best fit a non-linear model  $f(x, \mathbf{w})$  to a set of data points  $(x_i, y_i)$ .

# Non-Linear LS: the Levenberg-Marquardt algorithm.

## Goal

We want to find the parameters  $\mathbf{w}$  that best fit a non-linear model  $f(x, \mathbf{w})$  to a set of data points  $(x_i, y_i)$ .

## Mathematical Formulation

We aim to minimize the sum of squared residuals. Let's define the residuals as:

$$r_i(\mathbf{w}) = f(x, \mathbf{w}) - y_i$$

The total cost function  $f(\mathbf{w})$  to minimize is the sum of their squares:

$$f(\mathbf{w}) = \sum_{i=1}^m [r_i(\mathbf{w})]^2 = \mathbf{r}(\mathbf{w})^\top \mathbf{r}(\mathbf{w}) = \|\mathbf{r}(\mathbf{w})\|^2$$

# Non-Linear LS: the Levenberg-Marquardt algorithm.

## Goal

We want to find the parameters  $\mathbf{w}$  that best fit a non-linear model  $f(x, \mathbf{w})$  to a set of data points  $(x_i, y_i)$ .

## Mathematical Formulation

We aim to minimize the sum of squared residuals. Let's define the residuals as:

$$r_i(\mathbf{w}) = f(x, \mathbf{w}) - y_i$$

The total cost function  $f(\mathbf{w})$  to minimize is the sum of their squares:

$$f(\mathbf{w}) = \sum_{i=1}^m [r_i(\mathbf{w})]^2 = \mathbf{r}(\mathbf{w})^\top \mathbf{r}(\mathbf{w}) = \|\mathbf{r}(\mathbf{w})\|^2$$

The problem becomes:

$$\min_{\mathbf{w}} f(\mathbf{w})$$

# Two Approaches

## 1. Gradient Descent

- **Idea:** Move in the direction of the steepest descent of  $f(\mathbf{w})$ .
- **Update:**  
$$\mathbf{w}_{k+1} = \mathbf{w}_k - \gamma \nabla f(\mathbf{w}_k)$$
- **Pros:** Always converges to a minimum (stable).
- **Cons:** Very slow convergence, especially near the minimum. Only uses 1st-order info.

# Two Approaches

## 1. Gradient Descent

- **Idea:** Move in the direction of the steepest descent of  $f(\mathbf{w})$ .
- **Update:**  
$$\mathbf{w}_{k+1} = \mathbf{w}_k - \gamma \nabla f(\mathbf{w}_k)$$
- **Pros:** Always converges to a minimum (stable).
- **Cons:** Very slow convergence, especially near the minimum. Only uses 1st-order info.

## 2. Newton's Method

- **Idea:** Build a 2nd-order (quadratic) model of  $f(\mathbf{w})$  at each step and jump to its minimum.
- **Update:**  
$$\mathbf{w}_{k+1} = \mathbf{w}_k - H^{-1} \nabla f(\mathbf{w}_k)$$
- **Pros:** Very fast (quadratic) convergence.
- **Cons:** Requires computing the full Hessian matrix  $H$ , which is complex and computationally very expensive.

# The challenge

## Aims

We want to:

- have the speed of Newton's method;
- avoid the computation of the full Hessian.

# Derivation: the Gradient

We start with the goal of applying Newton's method, which requires the gradient  $\nabla f$  and the Hessian  $H = \nabla^2 f$ .

## 1. The Gradient $\nabla f$

The cost is  $f = \sum_i r_i^2$ . The  $j$ -th component of the gradient is:

$$\frac{\partial f}{\partial w_j} = \sum_{i=1}^m 2r_i \frac{\partial r_i}{\partial w_j}$$

Let  $J$  be the **Jacobian matrix** of the residuals  $\mathbf{r}$ , with entries  $J_{ij} = \partial r_i / \partial w_j$ .

In vector form, the gradient is:

$$\nabla f(\mathbf{w}) = 2J^\top \mathbf{r}(\mathbf{w})$$

# Derivation: the Hessian

## 2. The Full Hessian $H$

Differentiating the gradient again (with  $H_{jk} = \partial^2 f / \partial w_j \partial w_k$ ) gives:

$$H_{jk} = \sum_{i=1}^m 2 \left( \frac{\partial r_i}{\partial w_k} \frac{\partial r_i}{\partial w_j} + r_i \frac{\partial^2 r_i}{\partial w_j \partial w_k} \right)$$

In matrix form:

$$H = 2J^\top J + 2 \sum_{i=1}^m r_i \nabla^2 r_i$$

# The Gauss-Newton Approximation I

The full Hessian is:

$$H = \underbrace{2J^\top J}_{\text{First-order term}} + 2 \sum_{i=1}^m r_i \nabla^2 r_i \underbrace{\phantom{2 \sum_{i=1}^m r_i \nabla^2 r_i}}_{\text{Second-order term}}$$

## The Core Idea of Gauss-Newton

Approximate the Hessian by **dropping the second-order term**.

$$H \approx 2J^\top J$$

# The Gauss-Newton Approximation II

When is this a good approximation?

The second term  $2 \sum r_i \nabla^2 r_i$  is small if...

- **The residuals  $r_i$  are small.** This happens when the model is a good fit, i.e., we are **near the minimum**. This is why GN converges so fast near the solution.
- **The model is "almost linear".** If  $f(\mathbf{w})$  is close to linear, its second derivatives  $\nabla^2 r_i$  are close to zero.

# The Gauss-Newton Update Rule

- ① Start with the Newton step:

$$H\delta = -\nabla f$$

# The Gauss-Newton Update Rule

- ① **Start** with the Newton step:

$$H\delta = -\nabla f$$

- ② **Substitute** the gradient  $\nabla f = 2J^\top r$ :

$$H\delta = -2J^\top r$$

# The Gauss-Newton Update Rule

- ① **Start** with the Newton step:

$$H\delta = -\nabla f$$

- ② **Substitute** the gradient  $\nabla f = 2J^\top r$ :

$$H\delta = -2J^\top r$$

- ③ **Substitute** the approximate Hessian  $H \approx 2J^\top J$ :

$$(2J^\top J)\delta \approx -2J^\top r$$

# The Gauss-Newton Update Rule

- ① Start with the Newton step:

$$H\delta = -\nabla f$$

- ② Substitute the gradient  $\nabla f = 2J^\top r$ :

$$H\delta = -2J^\top r$$

- ③ Substitute the approximate Hessian  $H \approx 2J^\top J$ :

$$(2J^\top J)\delta \approx -2J^\top r$$

- ④ Simplify to get the **Gauss-Newton normal equations**:

$$(J^\top J)\delta = -J^\top r$$

# The Gauss-Newton Update Rule

- ① Start with the Newton step:

$$H\delta = -\nabla f$$

- ② Substitute the gradient  $\nabla f = 2J^\top r$ :

$$H\delta = -2J^\top r$$

- ③ Substitute the approximate Hessian  $H \approx 2J^\top J$ :

$$(2J^\top J)\delta \approx -2J^\top r$$

- ④ Simplify to get the **Gauss-Newton normal equations**:

$$(J^\top J)\delta = -J^\top r$$

## The Algorithm

At each step  $k$ , solve for the update  $\delta_k$  and apply it:

$$\mathbf{w}_{k+1} = \mathbf{w}_k + \delta_k$$

# Gauss-Newton: Advantages & Drawbacks

## Advantages

- **Fast Convergence:** If it converges, it does so quadratically (like full Newton) when near the minimum.
- **No Hessian Needed:** It only requires the Jacobian  $J$ , not the full, complex Hessian matrix  $H$ .
- $J^\top J$  is "cheaper" to compute than  $H$ .

# Gauss-Newton: Advantages & Drawbacks

## Advantages

- **Fast Convergence:** If it converges, it does so quadratically (like full Newton) when near the minimum.
- **No Hessian Needed:** It only requires the Jacobian  $J$ , not the full, complex Hessian matrix  $H$ .
- $J^\top J$  is "cheaper" to compute than  $H$ .

## Drawbacks

- **May Diverge:** If the initial guess is poor (far from minimum), the  $r_i$  are large, the Hessian approximation is bad, and the method can be unstable.
- **Singularity:** The matrix  $J^\top J$  may be singular or ill-conditioned, meaning we cannot solve for the step  $\delta$ . This happens if the Jacobian  $J$  does not have full column rank.

## The Dilemma

We want the **stability** of GD but the **speed** of GN.

# The Core Idea: A "Damped" Approach

## Levenberg-Marquardt (LM)

The LM algorithm adaptively interpolates between Gradient Descent and the Gauss-Newton method.

# The Core Idea: A "Damped" Approach

## Levenberg-Marquardt (LM)

The LM algorithm adaptively interpolates between Gradient Descent and the Gauss-Newton method.

It modifies the Gauss-Newton equation by adding a "damping" term  $\lambda$ :

## The LM Update Equation

$$(J^\top J + \lambda I)\delta = -J^\top r$$

# The Core Idea: A "Damped" Approach

## Levenberg-Marquardt (LM)

The LM algorithm adaptively interpolates between Gradient Descent and the Gauss-Newton method.

It modifies the Gauss-Newton equation by adding a "damping" term  $\lambda$ :

## The LM Update Equation

$$(J^\top J + \lambda I)\delta = -J^\top r$$

Where:

- $\delta$  is the update step ( $w_{k+1} = w_k + \delta$ ).
- $\lambda \geq 0$  is the **damping parameter**.
- $I$  is the identity matrix.

# The Core Idea: A "Damped" Approach

## Levenberg-Marquardt (LM)

The LM algorithm adaptively interpolates between Gradient Descent and the Gauss-Newton method.

It modifies the Gauss-Newton equation by adding a "damping" term  $\lambda$ :

## The LM Update Equation

$$(J^\top J + \lambda I)\delta = -J^\top r$$

Where:

- $\delta$  is the update step ( $w_{k+1} = w_k + \delta$ ).
- $\lambda \geq 0$  is the **damping parameter**.
- $I$  is the identity matrix.

How this solves the GN drawbacks:

By adding  $\lambda I$ , the matrix  $(J^\top J + \lambda I)$  is **always invertible** (for  $\lambda > 0$ ), even if  $J^\top J$  is singular.

# How the Damping Parameter $\lambda$ Works I

The LM update is:  $\delta = -(J^\top J + \lambda I)^{-1} J^\top r$

# How the Damping Parameter $\lambda$ Works I

The LM update is:  $\delta = -(J^\top J + \lambda I)^{-1} J^\top r$

Case 1:  $\lambda$  is Small (e.g.,  $\lambda \rightarrow 0$ )

$$(J^\top J + \lambda I) \approx J^\top J$$

The update becomes:

$$\delta \approx -(J^\top J)^{-1} J^\top r$$

**This is the Gauss-Newton method!**

- **Behavior:** Fast, but potentially unstable.
- **When:** Used when the current steps are good and reducing the error  $f$ .

# How the Damping Parameter $\lambda$ Works I

Case 2:  $\lambda$  is Large (e.g.,  $\lambda \rightarrow \infty$ )

For large  $\lambda$ , the  $J^\top J$  term is negligible:

$$(J^\top J + \lambda I) \approx \lambda I$$

The update becomes:

$$\delta \approx -(\lambda I)^{-1} J^\top r = -(1/\lambda) J^\top r$$

Recall  $\nabla f = 2J^\top r$ . So  $\delta \approx -(1/2\lambda)\nabla f$ .

**This is Gradient Descent!**

- **Behavior:** Slow, but stable. Takes small steps in the steepest direction.
- **When:** Used when steps are bad (increasing the error  $f$ ).

# The Algorithm: An Iterative Process

- ① Choose an initial guess  $w_0$  and an initial damping  $\lambda_0$ .

# The Algorithm: An Iterative Process

- ① Choose an initial guess  $\mathbf{w}_0$  and an initial damping  $\lambda_0$ .
- ② At each iteration  $k$ :
  - Compute residuals  $\mathbf{r}(\mathbf{w}_k)$  and Jacobian  $J(\mathbf{w}_k)$ .
  - Compute the current error  $f(\mathbf{w}_k)$ .
  - **Solve**  $(J^\top J + \lambda_k I)\delta_k = -J^\top \mathbf{r}$  for the step  $\delta_k$ .
  - Calculate a candidate new point:  $\mathbf{w}_{\text{new}} = \mathbf{w}_k + \delta_k$ .
  - Evaluate the new error  $f(\mathbf{w}_{\text{new}})$ .

# The Algorithm: An Iterative Process

- ① Choose an initial guess  $\mathbf{w}_0$  and an initial damping  $\lambda_0$ .
- ② At each iteration  $k$ :
  - Compute residuals  $\mathbf{r}(\mathbf{w}_k)$  and Jacobian  $J(\mathbf{w}_k)$ .
  - Compute the current error  $f(\mathbf{w}_k)$ .
  - Solve  $(J^\top J + \lambda_k I)\delta_k = -J^\top \mathbf{r}$  for the step  $\delta_k$ .
  - Calculate a candidate new point:  $\mathbf{w}_{\text{new}} = \mathbf{w}_k + \delta_k$ .
  - Evaluate the new error  $f(\mathbf{w}_{\text{new}})$ .
- ③ Update Strategy (The Core Logic):
  - If  $f(\mathbf{w}_{\text{new}}) < f(\mathbf{w}_k)$  (Good Step):
    - Accept the step:  $\mathbf{w}_{k+1} \leftarrow \mathbf{w}_{\text{new}}$ .
    - Trust the model more: Decrease  $\lambda$  (e.g.,  $\lambda_{k+1} = \lambda_k / \nu$ , with  $\nu > 1$ ). This moves towards Gauss-Newton.
  - If  $f(\mathbf{w}_{\text{new}}) \geq f(\mathbf{w}_k)$  (Bad Step):
    - Reject the step:  $\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k$ .
    - Trust the model less: Increase  $\lambda$  (e.g.,  $\lambda_{k+1} = \lambda_k \times \nu$ ). This moves towards Gradient Descent.

# The Algorithm: An Iterative Process

- ① Choose an initial guess  $\mathbf{w}_0$  and an initial damping  $\lambda_0$ .
- ② At each iteration  $k$ :
  - Compute residuals  $\mathbf{r}(\mathbf{w}_k)$  and Jacobian  $J(\mathbf{w}_k)$ .
  - Compute the current error  $f(\mathbf{w}_k)$ .
  - Solve  $(J^\top J + \lambda_k I)\delta_k = -J^\top \mathbf{r}$  for the step  $\delta_k$ .
  - Calculate a candidate new point:  $\mathbf{w}_{\text{new}} = \mathbf{w}_k + \delta_k$ .
  - Evaluate the new error  $f(\mathbf{w}_{\text{new}})$ .
- ③ Update Strategy (The Core Logic):
  - If  $f(\mathbf{w}_{\text{new}}) < f(\mathbf{w}_k)$  (Good Step):
    - Accept the step:  $\mathbf{w}_{k+1} \leftarrow \mathbf{w}_{\text{new}}$ .
    - Trust the model more: Decrease  $\lambda$  (e.g.,  $\lambda_{k+1} = \lambda_k / \nu$ , with  $\nu > 1$ ). This moves towards Gauss-Newton.
  - If  $f(\mathbf{w}_{\text{new}}) \geq f(\mathbf{w}_k)$  (Bad Step):
    - Reject the step:  $\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k$ .
    - Trust the model less: Increase  $\lambda$  (e.g.,  $\lambda_{k+1} = \lambda_k \times \nu$ ). This moves towards Gradient Descent.
- ④ Repeat until convergence (e.g., step size  $\delta$  is tiny, or error  $f$  stops decreasing).

# Advantages of Levenberg-Marquardt

## Robustness

It finds a solution even if the initial guess is far from the optimal minimum. By increasing  $\lambda$ , it can navigate "difficult" regions of the parameter space where Gauss-Newton would fail.

# Advantages of Levenberg-Marquardt

## Robustness

It finds a solution even if the initial guess is far from the optimal minimum. By increasing  $\lambda$ , it can navigate "difficult" regions of the parameter space where Gauss-Newton would fail.

## Efficiency

It combines the best of both worlds:

- Far from the minimum, it acts like stable Gradient Descent.
- Close to the minimum, it acts like fast-converging Gauss-Newton.

# Advantages of Levenberg-Marquardt

## Robustness

It finds a solution even if the initial guess is far from the optimal minimum. By increasing  $\lambda$ , it can navigate "difficult" regions of the parameter space where Gauss-Newton would fail.

## Efficiency

It combines the best of both worlds:

- Far from the minimum, it acts like stable Gradient Descent.
- Close to the minimum, it acts like fast-converging Gauss-Newton.

## Handles Ill-Conditioning

The damping term  $\lambda I$  **solves the main drawback of Gauss-Newton**.

The matrix  $(J^\top J + \lambda I)$  is always invertible (as long as  $\lambda > 0$ ).

## Computational Cost

- Requires computing (or approximating) the **Jacobian matrix**  $J$  at every iteration, which can be expensive for models with many parameters and data points.
- Requires solving an  $n \times n$  linear system (where  $n$  is # of parameters) at each iteration, an  $O(n^3)$  operation.

# Drawbacks and Considerations

## Computational Cost

- Requires computing (or approximating) the **Jacobian matrix**  $J$  at every iteration, which can be expensive for models with many parameters and data points.
- Requires solving an  $n \times n$  linear system (where  $n$  is # of parameters) at each iteration, an  $O(n^3)$  operation.

## Local Minimum

Like Gradient Descent and Gauss-Newton, LM is a **local optimizer**. It is only guaranteed to find a local minimum, not the global one. The final solution can depend heavily on the initial guess  $\mathbf{w}_0$ .

# Drawbacks and Considerations

## Computational Cost

- Requires computing (or approximating) the **Jacobian matrix**  $J$  at every iteration, which can be expensive for models with many parameters and data points.
- Requires solving an  $n \times n$  linear system (where  $n$  is # of parameters) at each iteration, an  $O(n^3)$  operation.

## Local Minimum

Like Gradient Descent and Gauss-Newton, LM is a **local optimizer**. It is only guaranteed to find a local minimum, not the global one. The final solution can depend heavily on the initial guess  $\mathbf{w}_0$ .

## Parameter Tuning

The performance can be sensitive to the initial value  $\lambda_0$  and the update factor  $\nu$ , though standard, "good-enough" values are widely used.

# Summary

- The Levenberg-Marquardt algorithm is the **de-facto standard** for solving non-linear least squares problems.

## Summary

- The Levenberg-Marquardt algorithm is the **de-facto standard** for solving non-linear least squares problems.
- It improves on Gauss-Newton by adding a **damping parameter**  $\lambda$  that makes the update step robust and invertible.

# Summary

- The Levenberg-Marquardt algorithm is the **de-facto standard** for solving non-linear least squares problems.
- It improves on Gauss-Newton by adding a **damping parameter**  $\lambda$  that makes the update step robust and invertible.
- It **adaptively interpolates** between:
  - Slow-but-stable **Gradient Descent** (when  $\lambda$  is large).
  - Fast-but-unstable **Gauss-Newton** (when  $\lambda$  is small).

# Summary

- The Levenberg-Marquardt algorithm is the **de-facto standard** for solving non-linear least squares problems.
- It improves on Gauss-Newton by adding a **damping parameter**  $\lambda$  that makes the update step robust and invertible.
- It **adaptively interpolates** between:
  - Slow-but-stable **Gradient Descent** (when  $\lambda$  is large).
  - Fast-but-unstable **Gauss-Newton** (when  $\lambda$  is small).
- It is robust, efficient, and handles ill-conditioned problems, but can be costly to compute and may get stuck in local minima.