

Matrix Completion Problem

The Netflix Challenge

The Problem:

- Netflix has millions of users
- Thousands of movies/shows
- Most users rate only a small fraction of content
- Goal: Predict missing ratings

Why it matters:

- Better recommendations
- Increased user engagement
- \$1M prize in 2009!



Key Insight: Low-Rank Structure

Intuition: User preferences are driven by a small number of latent factors

$$X \approx U \times V^T$$

Diagram illustrating matrix factorization:

- Matrix X (blue box) is approximately equal (\approx) to the product of matrices U (green box) and V^T (orange box).
- Matrix X has dimensions $n \times d$.
- Matrix U has dimensions $n \times r$ and is labeled "User features".
- Matrix V^T has dimensions $r \times d$ and is labeled "Movie features".

Examples of latent factors:

- Genre preferences (action, comedy, drama)
- Era preferences (classic vs. modern)
- Audience type (family, adult, teen)

Key assumption: $r \ll \min(n, d)$

Low-rank assumption

Low-rank assumption

For $n_1 \times n_2$ matrix M of rank r , assume the $\min(n_1, n_2) \gg r$.

- Why this assumption is needed?
- For simplicity, think about an $n \times n$ matrix M of rank r .
- It has $(2n - r)r$ degrees of freedom.
 - Degree of freedom is calculated by counting parameters in the SVD.
 - (The number of singular values)
 - + (degree of freedom of left singular vectors)
 - + (degree of freedom of right singular vectors)
$$= r + ((2n - r - 1)r)/2 + ((2n - r - 1)r)/2$$
 - Considerably smaller than n^2 .

Matrix Completion: Formal Setup

Given:

- A partially observed matrix with entries on index set
 $\Omega \subset \{1, \dots, n\} \times \{1, \dots, d\}$
- Observations: $\{X_{ij} : (i, j) \in \Omega\}$
- $|\Omega| = m$ (number of observed entries)

Goal: Recover the complete matrix $X \in \mathbb{R}^{n \times d}$ of rank r where

$$r \ll \min(n, d)$$

Assumption: The true underlying matrix has low rank, reflecting the presence of latent structure in the data.

Notation and Setup

Key objects:

- $X \in \mathbb{R}^{n \times d}$: True complete matrix (rank r)
- Ω : Set of observed indices
- $\mathcal{P}_\Omega(X)$: Projection onto observed entries

$$[\mathcal{P}_\Omega(X)]_{ij} = \begin{cases} X_{ij} & \text{if } (i,j) \in \Omega \\ 0 & \text{otherwise} \end{cases}$$

- Sampling fraction: $\rho = \frac{m}{nd} = \frac{|\Omega|}{nd}$

Challenge: This is an ill-posed problem without additional assumptions!

- Infinitely many matrices agree with observed entries
- Need to exploit low-rank structure

Ideal Estimator: Rank Minimization

Natural approach: Find the lowest-rank matrix consistent with observations

Rank Minimization Problem

$$\begin{aligned} & \underset{M \in \mathbb{R}^{n \times d}}{\text{minimize}} && \text{rank}(M) \\ & \text{subject to} && \mathcal{P}_\Omega(M) = \mathcal{P}_\Omega(X) \end{aligned}$$

Equivalently:

$$\underset{M \in \mathbb{R}^{n \times d}}{\text{minimize}} \quad \text{rank}(M) \quad \text{s.t.} \quad M_{ij} = X_{ij} \text{ for all } (i, j) \in \Omega$$


Problem: This is **NP-hard!**

- Rank function is non-convex
- Combinatorial in nature
- No efficient algorithms for general case


Non-convex vs Convex I

Visualizing Convex vs. Non-Convex Functions in 3D

Convex Function: $f(x, y) = x^2 + y^2$



Non-Convex Function with Multiple Minima



Non-convex vs Convex II

Convex.

- + Single global minimum
- + Simple to optimize; guaranteed convergence to global minimum
- Limited complexity: not suitable for real-world problems.

Non-convex.

- + Suitable for modeling landscapes related to real-world problems.
- Multiple local minima
- Convergence challenges: optimization methods can be stuck in local minima.

Nuclear Norm: The Convex Relaxation

Key idea: Replace rank with nuclear norm!

Nuclear Norm Definition

The **nuclear norm** (or trace norm) of a matrix M is:

$$\|M\|_* = \sum_{i=1}^{\min(n,d)} \sigma_i(M)$$

where $\sigma_1(M) \geq \sigma_2(M) \geq \dots \geq \sigma_{\min(n,d)}(M) \geq 0$ are the singular values of M .


Properties:

- Convex function
- Sum of singular values = trace of $\sqrt{M^T M}$
- ℓ_1 norm of the singular value vector


Rank vs Nuclear norm

$$M = [[\sqrt{4}x \ (x+y)], [(x+y) \ \sqrt{4}y]]$$

Rank of Matrix (Non-Convex)



Nuclear Norm (Convex Approximation)



Nuclear Norm Intuition

Analogy with sparse vectors:

Goal	Non-convex	Convex relaxation
Sparse vector	$\ x\ _0 = \#\{\text{non-zeros}\}$	$\ x\ _1 = \sum_i x_i $
Low-rank matrix	$\text{rank}(M)$	$\ M\ _* = \sum_i \sigma_i(M)$

Why it works:

- ℓ_1 norm promotes sparsity in vectors
- Nuclear norm promotes sparsity in singular values
- Sparsity in singular values = low rank!

Nuclear Norm Minimization

$$\begin{aligned} & \underset{M \in \mathbb{R}^{n \times d}}{\text{minimize}} && \|M\|_* \\ & \text{subject to} && \mathcal{P}_\Omega(M) = \mathcal{P}_\Omega(X) \end{aligned}$$

Advantages:

- Convex optimization problem
- Can be solved efficiently (though n, d large)
- Strong theoretical guarantees
- Exact recovery under mild conditions

When Does Nuclear Norm Work ? Theory I

Candes and Recht 2008:

Random Orthogonal Model

$$M = \sum_{k=1}^r \sigma_k u_k v_k^T$$

- $\{u_k\}_{k=1}^r$ is selected uniformly at random among all families of orthonormal vectors.
- The same for $\{v_k\}$.
- That is: M has an intrinsic low rank.

When Does Nuclear Norm Work ? Theory II

Theorem

Let $M \in \mathbb{R}^{n_1 \times n_2}$ be a rank r matrix sampled from the random orthogonal model. Let $n = \max(n_1, n_2)$. Suppose we observe m entries of M uniformly at random. Then, there are constants c and C such that if

$$m \geq Cn^{5/4}r \log n$$

then the minimizer to the Nuclear norm minimization is unique and equal to M with probability $1 - cn^{-3}$.

When Does Nuclear Norm Work ? Theory III

Implications:

- Under the hypothesis of the Theorem, there is a unique low-rank matrix which is consistent with the observed entries.
- This matrix can be recovered by a convex optimization algorithm.
- So the nuclear norm relaxation is formally equivalent to the NP hard problem.

Computational Challenge

Problem: Even though convex, directly solving nuclear norm minimization is expensive for large matrices.

Solution: Singular Value Thresholding (SVT) Algorithm

- Iterative algorithm
- Exploits problem structure
- Only requires SVD of sparse matrices
- Efficient for large-scale problems

Key Ingredient: Soft-Thresholding of Singular Values

Scalar soft-thresholding:

$$\mathcal{S}_\tau(x) = \text{sign}(x) \cdot \max(|x| - \tau, 0) = \begin{cases} x - \tau & \text{if } x > \tau \\ 0 & \text{if } |x| \leq \tau \\ x + \tau & \text{if } x < -\tau \end{cases}$$

Matrix soft-thresholding: Let $M = U\Sigma V^T$ be the SVD with $\Sigma = \text{diag}(\sigma_1, \dots, \sigma_r)$


Singular Value Thresholding Operator

$$\mathcal{D}_\tau(M) = US_\tau(\Sigma)V^T = \sum_{i=1}^r \mathcal{S}_\tau(\sigma_i) u_i v_i^T$$

where $\mathcal{S}_\tau(\Sigma) = \text{diag}(\mathcal{S}_\tau(\sigma_1), \dots, \mathcal{S}_\tau(\sigma_r))$

Interpretation: Shrinks singular values by τ , sets small ones to zero

Soft-Thresholding Visualization



Singular values below τ are set to zero

Effect: Promotes low rank by eliminating small singular values

Why Soft-Thresholding?

Connection between \mathcal{D}_τ and $\tau \|\cdot\|_*$

The soft-thresholding operator \mathcal{D}_τ can be written also as:

$$\mathcal{D}_\tau(Y) = \arg \min_M \left\{ \frac{1}{2} \|M - Y\|_F^2 + \tau \|M\|_* \right\}$$

This means:

- Soft-thresholding solves a regularized least-squares problem
- Balances fidelity to Y with nuclear norm penalty
- Closed-form solution via SVD!

Computation:

- ① Compute SVD: $Y = U\Sigma V^T$
- ② Threshold: $\Sigma' = \mathcal{S}_\tau(\Sigma)$
- ③ Reconstruct: $\mathcal{D}_\tau(Y) = U\Sigma'V^T$

SVT Algorithm: Setup

Reformulation: Consider the unconstrained problem

$$\underset{M}{\text{minimize}} \quad \tau \|M\|_* + \frac{1}{2} \|\mathcal{P}_\Omega(M) - \mathcal{P}_\Omega(X)\|_F^2$$

for appropriate choice of $\tau > 0$.

Equivalent constrained form:

$$\begin{aligned} & \underset{M}{\text{minimize}} \quad \tau \|M\|_* \\ & \text{subject to} \quad \mathcal{P}_\Omega(M) = \mathcal{P}_\Omega(X) \end{aligned}$$

The SVT algorithm solves this using an iterative projection method.

SVT Algorithm: The Iteration

Singular Value Thresholding Algorithm

Input: Observations $\mathcal{P}_\Omega(X)$, parameter $\tau > 0$, step size $\delta \in (0, 2)$

Initialize: $Y_0 = c_0 \delta P_\Omega(X)$, $k = 0$

Repeat until convergence:

- ① Compute: $M_k = \mathcal{D}_\tau(Y_k)$
- ② Update: $Y_{k+1} = Y_k + \delta \mathcal{P}_\Omega(X - M_k)$
- ③ $k \leftarrow k + 1$

Output: M_k

Interpretation:

- Step 1: Shrink singular values (promote low rank)
- Step 2: Project back to agreement with observations
- Alternates between rank minimization and data fidelity

SVT Algorithm: Detailed Steps

Step-by-step breakdown:

Step 1: Singular value thresholding

- Compute SVD: $Y_k = U_k \Sigma_k V_k^T$
- Apply soft-threshold: $M_k = U_k \mathcal{S}_\tau(\Sigma_k) V_k^T$
- This gives a low-rank approximation to Y_k

Step 2: Gradient update

- Compute residual: $R_k = X - M_k$ on Ω
- Update: $Y_{k+1} = Y_k + \delta \mathcal{P}_\Omega(R_k)$
- Move in direction that reduces constraint violation
- δ controls step size (typically $\delta = 1.2$)

Key: Only need to compute SVD once per iteration, and only for top singular vectors!

SVT: Convergence Properties

Theoretical guarantees:

Theorem (Cai, Candès, Shen 2010)

For $\delta \in (0, 2)$, the SVT algorithm converges to the optimal solution of the nuclear norm minimization problem.

Convergence rate:

- Linear convergence: $\|M_k - M^*\|_F \leq C\rho^k$ for some $\rho < 1$
- Rate depends on δ and problem conditioning
- Best choice: $\delta \approx 1.2$ to 1.5 (empirically)

Stopping criterion:

$$\frac{\|\mathcal{P}_\Omega(M_k) - \mathcal{P}_\Omega(X)\|_F}{\|\mathcal{P}_\Omega(X)\|_F} < \epsilon$$

for desired tolerance ϵ (e.g., $\epsilon = 10^{-4}$).

SVT: Computational Complexity

Per iteration cost:

① **SVD computation:** $O(r^2(n + d))$ using partial SVD

- Only need top r singular vectors
- Can use iterative methods (Lanczos, power method)
- Much cheaper than full SVD: $O(\min(n^2d, nd^2))$

② **Projection update:** $O(m)$ where $m = |\Omega|$

- Only update observed entries
- Sparse operation

Overall complexity:

$$O(K \cdot r^2(n + d))$$

where K is the number of iterations.

Scalability: Efficient for large n, d when r is small!

SVT: Key Properties

Advantages:

- **Simple implementation:** Only requires SVD and basic operations
- **Memory efficient:** Can work with sparse representations
- **Scalable:** Handles large matrices when rank is small
- **Guaranteed convergence:** Provably converges to optimal solution
- **Flexible:** Easy to add constraints or regularization

Practical considerations:

- Choice of τ : Cross-validation or theoretical bounds
- Warm start: Initialize with previous solution
- Adaptive δ : Can adjust step size during iteration
- Early stopping: Don't need exact convergence

SVT: Parameter Selection

Choosing τ :

① Theoretical choice:

$$\tau = \gamma \sqrt{nd}$$

where $\gamma \approx 5$ to 10 (rule of thumb)

② Cross-validation:

- Hold out some observed entries
- Select τ minimizing validation error
- Grid search over $\tau \in [1, 10\sqrt{nd}]$

③ Adaptive:

$$\tau_k = \tau_0 \cdot \alpha^k$$

Decrease τ over iterations (continuation method)

Choosing δ : Typical values $\delta \in [1.0, 1.5]$

- Larger δ : Faster initial progress
- Smaller δ : More stable convergence

SVT: Practical Algorithm

Algorithm 1 Practical SVT Implementation

Require: $\mathcal{P}_\Omega(X)$, τ , δ , ϵ , k_{\max}

- 1: Initialize: $Y \leftarrow c_0 \delta P_\Omega(X)$, $k \leftarrow 0$
- 2: **while** $k < k_{\max}$ **do**
- 3: $[U, \Sigma, V] \leftarrow \text{partialSVD}(Y, r_k)$ {Top r_k components}
- 4: $\Sigma' \leftarrow \max(\Sigma - \tau, 0)$ {Soft-threshold}
- 5: $M \leftarrow U\Sigma'V^T$
- 6: $R \leftarrow \mathcal{P}_\Omega(X - M)$ {Residual on Ω }
- 7: **if** $\|R\|_F / \|\mathcal{P}_\Omega(X)\|_F < \epsilon$ **then**
- 8: **break** {Converged}
- 9: **end if**
- 10: $Y \leftarrow Y + \delta R$
- 11: $r_{k+1} \leftarrow \text{rank}(M) + 1$ {Adaptive rank}
- 12: $k \leftarrow k + 1$
- 13: **end while**
- 14: **return** M

Alternative Algorithms

Other popular methods:

Algorithm	Key Idea
OptSpace	Low-rank factorization on manifold
SOFT-IMPUTE	Iterative soft-thresholding with warm starts
FPCA	Fixed point iteration with continuation
APGL	Accelerated proximal gradient (Nesterov)
LMaFit	Low-rank matrix fitting (non-convex)
ALS	Alternating least squares (factorization)

Comparison with SVT:

- SVT: Simple, guaranteed convergence, may be slower
- OptSpace: Very fast, but needs good initialization
- SOFT-IMPUTE: Similar to SVT, good practical performance
- Non-convex methods: Faster but no guarantees

SVT vs. Factorization Methods

SVT (Convex)

- + Global optimum
- + Theoretical guarantees
- + No initialization needed
- Slower convergence
- SVD per iteration
- Memory for full M

Factorization (Non-convex)

- + Very fast
- + Memory efficient
- + No SVD needed
- Local minima
- Initialization matters
- Rank must be specified

Recommendation:

- Small/medium problems: Use SVT
- Large-scale: Consider factorization or SOFT-IMPUTE
- When guarantees matter: Use SVT

Applications of Matrix Completion

① Computer Vision

- Image inpainting
- Video completion
- Background subtraction

② System Identification

- Missing sensor data
- Network topology inference
- Wireless sensor networks

③ Bioinformatics

- Gene expression data
- Protein interaction networks
- Drug response prediction

Applications (continued)

④ Quantum State Tomography

- Reconstruct quantum states
- Reduce measurement complexity


⑤ Collaborative Filtering

- Product recommendations
- Social network analysis
- Ad targeting

⑥ Financial Data

- Portfolio optimization with missing returns
- Risk assessment
- Market prediction

Example: Image Inpainting



Key insight: Natural images have low-rank structure when properly arranged (e.g., patch-based representations).

Summary: Key Takeaways

- ① **Problem:** Complete partially observed low-rank matrix
 - Motivated by recommendation systems
 - Assumes rank $r \ll \min(n, d)$
- ② **Approach:** Convex relaxation
 - Replace rank with nuclear norm
 - Provable recovery guarantees
- ③ **Algorithm:** Singular Value Thresholding
 - Iterative soft-thresholding of singular values
 - Computationally efficient for low-rank matrices
 - Guaranteed convergence
- ④ **Applications:** Widespread in machine learning and signal processing