


# Activation Functions

# 1. Binary Step Function - Figure

Binary Step Function



$$f(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ 0 & \text{if } x < 0 \end{cases}$$

# 1. Binary Step Function - Analysis

## Properties

- Output is either 0 or 1 ("fires" or "doesn't fire").
- A non-linear function.


## Drawbacks

- **Zero Gradient:** The derivative is 0 everywhere except at  $x = 0$  (where it's undefined).
- This means **backpropagation cannot work** as no gradient flows back to update weights.
- Cannot be used for multi-class classification.

## Use

- Rarely used in modern networks.
- Historically important (e.g., in the original **Perceptron**).

## 2. Linear Function - Figure



$$f(x) = x$$

## 2. Linear Function - Analysis

### Properties

**Identity function:** Output is equal to the input.

Derivative is constant ( $f'(x) = 1$ ).

### Drawbacks


- **No Non-linearity:** Using this in hidden layers is pointless.
- A stack of linear layers is mathematically equivalent to a **single linear layer**.
- The network cannot learn complex, non-linear patterns.

### Use

- Typically only used in the **output layer** for **regression tasks** (where the output value can be any real number).

### 3. Sigmoid / Logistic - Figure

Sigmoid Function



$$f(x) = \sigma(x) = \frac{1}{1+e^{-x}}$$

### 3. Sigmoid / Logistic - Analysis

#### Properties

- S-shaped curve, **Range:** (0, 1).
- Differentiable and smooth.

#### Advantages

- Output can be interpreted as a **probability**.
- Smooth gradient.


#### Drawbacks

- **Vanishing Gradients:** For large positive or negative  $x$ , the gradient  $\approx 0$ .
- **Not zero-centered:** Output is always positive. This can make training unstable (all gradients for weights in a layer will have the same sign).

#### Use

- **Output layer for binary classification.**
- Rarely used in hidden layers today (Tanh or ReLU are preferred).

#### 4. Tanh (Hyperbolic Tangent) - Figure



$$f(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

## 4. Tanh (Hyperbolic Tangent) - Analysis

### Properties

- S-shaped, like a "scaled" sigmoid.
- **Range:**  $(-1, 1)$ .
- Differentiable and smooth.

### Advantages

- **Zero-centered:** Output mean is near 0. This helps center the data for the next layer and makes training easier and more stable than Sigmoid.


### Drawbacks

- Still suffers from the **vanishing gradient problem** at the extremes (though the gradient is steeper than Sigmoid).

### Use

- Commonly used in **hidden layers** of standard feed-forward networks.
- Very common in **Recurrent Neural Networks (RNNs)**.
- Generally preferred over Sigmoid for hidden layers.

## 5. ReLU (Rectified Linear Unit) - Figure



$$f(x) = \max(0, x)$$

## 5. ReLU (Rectified Linear Unit) - Analysis

### Properties

- Linear for  $x > 0$ , zero for  $x \leq 0$ .
- Non-saturating (for  $x > 0$ ).

### Advantages

- **Computationally efficient:** Very simple max operation.
- **Solves vanishing gradient** for positive inputs.
- **Sparsity:** Induces sparsity as many neurons output 0.

### Drawbacks


- **The "Dying ReLU" Problem:** If a neuron's input is always negative (e.g., due to a large negative bias), it gets "stuck" outputting 0. Its gradient will always be 0, and it will never learn again.
- **Not zero-centered.**

### Use

- The **most common** activation function for **hidden layers** in **CNNs** and **DNNs**. It is the default choice.

## 6. Leaky ReLU - Figure

Leaky ReLU Function



$$f(x) = \begin{cases} x & \text{if } x \geq 0 \\ 0.01x & \text{if } x < 0 \end{cases}$$

## 6. Leaky ReLU - Analysis

### Properties

- An improved version of ReLU.
- Introduces a small, non-zero slope for negative inputs (e.g.,  $\alpha = 0.01$ ).

### Advantages

- **Solves the "Dying ReLU" Problem:** The gradient is never 0, so neurons can always recover and continue learning.
- Retains the efficiency and non-saturation benefits of ReLU.

### Drawbacks


- Results are not always consistent; sometimes it works better than ReLU, sometimes not.
- The slope  $\alpha$  (0.01) is a fixed hyperparameter that must be chosen.

### Use

- A common alternative to ReLU, especially if "Dying ReLU" is a suspected problem.

## 7. Parameterized ReLU (PReLU) - Figure

PReLU Function (example  $\alpha = 0.2$ )



$$f(x) = \begin{cases} x & \text{if } x \geq 0 \\ \alpha x & \text{if } x < 0 \end{cases}$$

## 7. Parameterized ReLU (PReLU) - Analysis

### Properties

- A variation of Leaky ReLU where  $\alpha$  is a **learnable parameter**.

### Advantages

- The network **learns the best slope**  $\alpha$  during backpropagation.
- Can adapt to the data better than Leaky ReLU.
- Solves the "Dying ReLU" problem.

### Drawbacks


- Can be prone to **overfitting** on smaller datasets, as it adds more parameters to the model.
- Increases model complexity.

### Use

- Used in more complex tasks where the extra learnable parameter might capture useful information.

## 8. Exponential Linear Unit (ELU) - Figure

ELU Function (example  $\alpha = 1.0$ )



$$f(x) = \begin{cases} x & \text{if } x \geq 0 \\ \alpha(e^x - 1) & \text{if } x < 0 \end{cases}$$

## 8. Exponential Linear Unit (ELU) - Analysis

### Properties

- Smooth curve for negative inputs, unlike the sharp corner in ReLU.
- Becomes  $-\alpha$  for large negative  $x$ .

### Advantages

- Solves "Dying ReLU" problem.
- Near zero-centered output, which improves on standard ReLU.
- The smooth curve can help the optimizer converge faster.


### Drawbacks

- More computationally expensive than ReLU due to the exponential function ( $e^x$ ).
- $\alpha$  is a hyperparameter to tune.

### Use

- A strong alternative to ReLU, especially when faster convergence and better accuracy are needed (and compute cost is less of an issue).

## 9. Scaled ELU (SELU) - Figure



$$f(x) = \lambda \begin{cases} x & \text{if } x > 0 \\ \alpha(e^x - 1) & \text{if } x \leq 0 \end{cases}$$

## 9. Scaled ELU (SELU) - Analysis

### Properties

- A specific, scaled version of ELU.
- The values for  $\lambda$  and  $\alpha$  are pre-defined constants ( $\lambda \approx 1.0507$ ,  $\alpha \approx 1.6732$ ).

### Advantages

- **Self-Normalizing:** If weights are initialized correctly (Lecun Normal), SELU can push neuron activations towards a mean of 0 and std. dev. of 1.
- This "self-normalizing" property acts like internal batch normalization, **preventing vanishing/exploding gradients**.

### Drawbacks

- Requires strict conditions: **Lecun Normal weight initialization** and **AlphaDropout** (not standard dropout).
- Not compatible with all network architectures (e.g., CNNs).

### Use

- Designed for **deep Feed-Forward Networks (FNNs)**.


## 10. Softmax Function - Concept

### Definition

This function is different: it's applied to a **vector** of inputs (the final layer's logits), not a single number. It converts a vector of raw scores into a **probability distribution**.

$$\text{Softmax}(x_i) = \frac{e^{x_i}}{\sum_{j=1}^K e^{x_j}}$$

Where  $x_i$  is the  $i^{th}$  element of the input vector and  $K$  is the total number of classes.



**Sum = 1.0**

## 10. Softmax Function - Analysis

### Properties & Advantages

- **Range:** Each output element is in  $(0, 1)$ .
- **Sum to 1:** All elements in the output vector sum to 1.
- **Interpretable:** The output can be seen as the model's confidence or probability for each class.
- Differentiable.


### Drawbacks

- Can be computationally expensive (many exponentials).
- Can still suffer from saturation if input values (logits) are very large or very small.

### Typical Use

- **Exclusively** used in the **output layer** for **multi-class classification** problems.

## 11. Swish - Figure



$$f(x) = x \cdot \sigma(x)$$

## 11. Swish - Analysis

### Properties

- $f(x) = x \cdot \text{Sigmoid}(x)$ .
- **Non-monotonic:** It dips slightly for negative values.
- Smooth and differentiable everywhere.
- Unbounded above, bounded below.

### Advantages

- Tends to outperform ReLU on deeper networks.
- The non-monotonicity is claimed to help with learning.
- Avoids the "Dying ReLU" problem.


### Drawbacks

- **Computationally expensive** (due to Sigmoid).

### Use

- Developed by Google and used in their EfficientNet models [▶ Link](#).
- A strong, modern alternative to ReLU for deep networks.

## 12. Mish Function - Figure



$$f(x) = x \cdot \tanh(\text{softplus}(x))$$

## 12. Mish Function - Analysis

### Properties

- $\text{softplus}(x) = \ln(1 + e^x)$ .
- Similar shape to Swish (non-monotonic, smooth).
- Unbounded above, bounded below.

### Advantages

- Reported to provide even better performance than Swish and ReLU in many deep learning tasks.
- Smoothness helps with optimization and generalization.

### Drawbacks

- Even **more computationally expensive** than Swish due to the combination of  $\tanh$ ,  $\ln$ , and  $e^x$ .

### Use

- A state-of-the-art activation function used in object detection (e.g., YOLOv4) and other complex vision tasks.

# How to Choose the Right Function?

## General Rules of Thumb

- **Start with ReLU:** It's fast, efficient, and works well for most problems (especially CNNs). Be mindful of "Dying ReLU".
- **If ReLU fails:** Try **Leaky ReLU**, **ELU**, or **PReLU** to fix the "dying neuron" problem.
- **For deep FNNs:** Consider **SELU** (with Lecun initialization) to leverage self-normalization.
- **For RNNs:** **Tanh** or **Sigmoid** are still standard choices.
- **For state-of-the-art:** If compute is not a bottleneck, experiment with **Swish** or **Mish**.

# How to Choose the Right Function?

## General Rules of Thumb

- **Start with ReLU:** It's fast, efficient, and works well for most problems (especially CNNs). Be mindful of "Dying ReLU".
- **If ReLU fails:** Try **Leaky ReLU**, **ELU**, or **PReLU** to fix the "dying neuron" problem.
- **For deep FNNs:** Consider **SELU** (with Lecun initialization) to leverage self-normalization.
- **For RNNs:** **Tanh** or **Sigmoid** are still standard choices.
- **For state-of-the-art:** If compute is not a bottleneck, experiment with **Swish** or **Mish**.

## For the Output Layer

The choice is dictated by the problem:

- **Binary Classification:** Sigmoid
- **Multi-Class Classification:** Softmax
- **Regression (any value):** Linear